Stochastic Discrete Local Volatility

Abstract

We present a fully arbitrage-free option market model using the SANOS [6] framework based in ideas of [19].

The model is the product of a risk-neutral discrete state martingale defined on a fixed grid with stochastic "discrete local volatilities", and a simple SANOS-type Black & Scholes smoothing process which allows smooth interpolation of option prices between grid points.

We train the model on S&P 500 Index options data.

Abstract

We study option market models on a fixed lattice of cash strikes and expiry dates whose state is the vector of gridded risk-neutral densities together with the discrete local variances. Positivity of the local variances is equivalent to absence of static arbitrage including calendar spreads, so the admissible set is an open orthant and $\lambda = \log \ell$ ranges over a Euclidean space: a diffusion on the code-book is statically arbitrage-free by construction.

Two results follow from the choice of chart. First, the spot is a read-off, $S_t = x^\top q_t^j$ at every expiry, an identity of the chart holding at all t with no limit taken. At the front expiry we exhibit a family of front volatilities under which the density collapses onto a single node, so the spot lands on a strike and every gridded contract settles at its exact payoff. That pinning is a property of the chosen volatility and not of the chart, which represents every terminal spot in the grid range; in [19] it follows from the chart whatever the coefficients. Second, the model has a unique global strong solution on $[0, T_1]$, up to and including the first expiry, for every bounded and locally Lipschitz volatility structure and any number of driving factors, with no damping, no frozen wing, and no design rule. The only requirement is that the front volatility be integrable in time, which the natural collapse rate $(T_1 - t)^{-1/2}$ satisfies. The quantity to be controlled is $\eta = \psi/q$, the density volatility relative to the density; it is bounded by the grid alone. Wissel [19] must assume a bound of this kind; here it is a theorem.

The no-arbitrage drift restriction itself is not new. In the continuum it is [9], and on this lattice it is [19], whose local implied variance coincides with the discrete local variance used here as an algebraic identity on the grid. What the density chart adds is that the drift becomes pointwise, $\alpha_i^j = -\sum_f \eta_i^{j,f} \beta_i^{j,f}$, with no matrix inversion and no rank condition, the apparent over-determination being resolved by conservation of mass and of the spot.

1 Introduction

This article presents an arbitrage-free market model for for smooth surfaces of European option prices.

The core architecture follows the SANOS approach [6] which models the stock price as the product of a time-and-state discrete process X and a smoothing process Y, in our case Black & Scholes. The former is used to encode the option market using a discrete form of local volatilities as code book. That does not mean that the model is a local volatility model – we

merely use the (discrete) local volatilities as a convenient representation of the option surface. Compared to using implied volatilities the advantage is that mere positivity of the local volatilities is sufficient to ensure the the option market is free of static arbitrage at all times. We have shown in [6] that this representation can be fitted to real life market data very efficiently using linear programming.

In [7] we have shown how to estimate the dynamics such (discrete) local volatilities of such a model under the statistical measure $\mathbb{P}$ which fit very well to real-life observed option bid/asks.

The key contribution in this paper is showing how to construct a corresponding risk-neutral measure $\mathbb{Q}$ or, equivalently, the risk-neutral drift of the model. Our approach is closely related to the seminal but unpublished work [19].

An option market model reverses the usual direction of construction: instead of specifying a process for the underlying and deriving option prices from it, one models the code-book of prices directly, so that a calibration is a change of coordinates rather than an inverse problem. Two conditions must then be maintained at once. The state must remain free of *static* arbitrage, and its dynamics must admit one equivalent local martingale measure common to the underlying and to every traded option. The second is the option analogue of an HJM drift restriction: the volatility structure may be specified freely, and the drift is then determined. The lattice here is *finite*, finitely many strikes and finitely many expiry dates, which is what is traded.

What is already known. Schweizer and Wissel treat the problem in the continuum in two companion papers, one in each direction of the surface: [15] varies maturity at a single strike, and [14] varies strike at a single maturity. In each the code-book is chosen so that its positivity encodes absence of arbitrage, the drift restriction is derived, and existence is proved. The multi-strike multi-maturity case, which is ours, is stated as open in [14], with the expectation that a new parametrisation will be needed. Carmona and Nadtochiy [9] give a continuum local-volatility code-book with a drift restriction, but no existence theorem. On a finite lattice, and answering the multi-maturity question, stands Wissel [19]; his construction is the one this paper is closest to, and it is described separately below.

Modern approaches. The question was reopened from the data side, by the scarcity of option-market histories for training and evaluating hedging strategies [5]. Cohen, Reisinger and Wang [10] build neural stochastic differential equations on a floating lattice and penalise a least-squares residual in place of the drift restriction. [16] simulate an equity option book with a generative adversarial network whose state is the discrete local variance grid compressed by principal components, so that static arbitrage is absent by positivity of that grid; that is the code-book choice this paper inherits, and [17] extend it to a multi-asset universe. [7] decode statically arbitrage-free surfaces from a low-dimensional code, fit a linear state process to it, and test for dynamic arbitrage after the fact, reporting it in 12.5% of states at zero transaction cost.

None of the four imposes a dynamic condition, and it is where this paper starts: the cross-sectional constraints govern each date's surface and its terminal payoff cone, and leave the transition of a fixed contract between two dates unconstrained. The choice of grid is part of the same point. All four index their nodes relative to the spot rather than in fixed cash terms, and on such a grid a node is not one contract through time, so there is no single instrument whose price is required to be a martingale and nothing for an HJM-type restriction to constrain. Fixing the strikes in cash is what buys the restriction here.

The closest prior work. Wissel [19] works on a finite lattice of fixed cash strikes and fixed expiry dates, parametrises it by a price level together with positive local implied volatilities, derives the no-arbitrage drift restrictions, and proves constructive strong existence and uniqueness for a class of coefficients. Behind the front expiry his local implied variance is *exactly* the backward discrete local variance used here, an algebraic identity on the grid, and his drift restriction is, after a change of measure, exactly the one we state in Theorem 3.3. We claim no priority for it; see Appendix A.2, which gives the identification in full. That working paper appears never to have been published in archival form.

What this paper does

We keep the same lattice and change the chart: the state is the vector of gridded *densities* together with the discrete local variances, rather than prices together with quantiles of digitals. Three consequences follow, and they are the contributions.

1. **The spot is a read-off, and it settles exactly.** The spot is $S_t = x^\top q_t^j$ at every expiry j (Proposition 3.4), an identity of the chart holding at all t. At the front expiry Proposition 4.3 exhibits a family of front volatilities under which q^1 collapses onto a single node, so the spot lands on a strike and every gridded contract settles at its exact payoff (Section 4.2). That pinning is a property of the volatility chosen and not of the chart: by Remark 4.4 the terminal states the chart admits carry every spot in the grid range, whereas in [19] the pinning follows from the chart whatever the coefficients (Remark A.1). We do not construct dynamics reaching an off-lattice terminal spot; that is left as future work.

2. **Existence, with no condition on the wings.** The model has a unique global strong solution on $[0, T_1]$, positive by construction, for every bounded and locally Lipschitz volatility structure b and any number of driving factors. The quantity that must be controlled is $\eta = \psi/q$, the density volatility relative to the density, and it is bounded by the grid alone: Theorem 5.6 gives $\Theta_j \le N + 1$ uniformly in the state and in the volatility structure. *Wissel [19] must assume a bound of this kind; here it is a theorem.* The front expiry costs one hypothesis, and it is a mild one: η^1 must be integrable in t, not bounded, because η^j is affine in η^1 and an Itô equation needs an integrable drift rather than a bounded one. The front volatility used here diverges at rate $(T_1 - t)^{-p}$, and the two demands made of p leave a window: $p < 1$ is what Assumption 5.7 asks, while $p \ge \frac{1}{2}$ is what makes the front density collapse, so that every gridded contract settles at its exact payoff (Proposition 4.3).

3. **The drift is pointwise, and the count works out.** In density coordinates the drift equation is $N + 2$ conditions on N unknowns, and is consistent because the left null space of the geometry matrix is exactly $\mathrm{span}\{\mathbf{1}, x\}$, conservation of mass and of the spot (Lemma 2.3). The state-independent factor then cancels from both sides and the drift is available nodewise, $\alpha_i^j = -\sum_f \eta_i^{j,f} \beta_i^{j,f}$, with no matrix inversion and no rank condition.

Scope, and what is not claimed

The over-determination of the third contribution is not intrinsic: it is created by the choice of density coordinates and removed by the same choice, and in price coordinates the system is square and the question does not arise.

The existence statement is for a fixed finite grid. The constant is $(M-1)(N+1)$ and depends on the grid. The one hypothesis, Assumption 5.7, is a condition on the front volatility and not on the state, and it is sufficient rather than necessary.

2 The code-book

We are assuming deterministic forwards and discount factors which is equivalent to assuming that spot is driftless and discount factors are one, c.f. [4]. Assume that spot is S_0 at time $t = 0$.

We follow the [6] architecture for a time-homogeneous fixed strike grid: Fix model expiries $0 =: T_0 < T_1 < \cdots < T_M$ with $\Delta_j := T_j - T_{j-1}$ and $T_0 := 0$, and a single strike grid $0 \le x^0 < x^1 < \cdots < S_0 < \ldots x^N < x^{N+1}$, held fixed in cash terms and common to all expiries. Let $q = (q^{j,i})$ with $q^j \in \mathbb{R}^{N+2}$ be a discrete risk-neutral density in the sense that

1. $q^{j,i} \ge 0$
2. $\sum_{i=0}^{N+1} q^{j,i} = 1$
3. $\sum_{i=0}^{N+1} x_i q^{j,i} =: S_0$.

We denote by $\Sigma_t^{j,i}$ the discrete local volatility squared

on that grid as seen at calendar time t, and let $\ell_t^j \in (0,\infty)^N$ be the discrete local variance on the interior nodes, with $\lambda = \log \ell$ throughout. Prices are the linear functional $C_t(x_i, T_j) = \pi_i' q_t^j$ with π_i the discretised payoff $(x - x_i)^+$.

If it's the same for everyone I'd align the notation with DYSANOS, e.g. use $\hat{T}$ and $\hat{K}$ for relative model coordinates, and T, K for the market. I'd generally use the same terms. I would also write everything in discrete local vol-squared and specifically add the "dt" term because otherwise the numbers have quite different dimensionality.

Propagating densities rather than prices, in Fokker–Planck form, is what makes the systems below tridiagonal. With the half-weights

$$\omega_k^j := \frac{x_k^2\,\Delta_j}{x_{k+1} - x_{k-1}}, \qquad \iota_k^{j,-} := \frac{\omega_k^j}{x_k - x_{k-1}}, \qquad \iota_k^{j,+} := \frac{\omega_k^j}{x_{k+1} - x_k}, \qquad k = 1, \ldots, N, \tag{1}$$

the recursion is $A_j(\ell^j)\, q^j = q^{j-1}$, where $A_j(\ell)$ is tridiagonal with rows

$$\big(A_j(\ell)q\big)_i = \begin{cases} q_0 - \ell_1 \iota_1^{j,-} q_1, & i = 0, \\ -\ell_{i-1}\iota_{i-1}^{j,+} q_{i-1} + \big(1 + \ell_i(\iota_i^{j,-} + \iota_i^{j,+})\big) q_i - \ell_{i+1}\iota_{i+1}^{j,-} q_{i+1}, & 1 \le i \le N, \\ -\ell_N \iota_N^{j,+} q_N + q_{N+1}, & i = N+1. \end{cases} \tag{2}$$

Up to notation this is the implicit forward operator of [8]. Note that in (2) the entry multiplying $q_{i\pm1}$ carries $\ell_{i\pm1}$ and not ℓ_i.

Because (2) is affine in ℓ we may write

$$A_j(\ell)\, q \;=\; q + G_j(q)\, \ell, \qquad G_j(q) \in \mathbb{R}^{(N+2)\times N}, \tag{3}$$

so that $G_j(q)$ collects the terms carrying ℓ. The map G_j is linear in q and $G_j(q)\ell$ is linear in each argument separately; we use both readings. The following elementary observation is used constantly.

Lemma 2.1 (Factorisation). *Column k of $G_j(q)$ involves q only through q_k. Consequently*

$$G_j(q) \;=\; \widehat{G}_j \,\mathrm{diag}(q_1, \ldots, q_N), \tag{4}$$

where $\widehat{G}_j \in \mathbb{R}^{(N+2)\times N}$ has k-th column $-\iota_k^{j,-} e_{k-1} + (\iota_k^{j,-} + \iota_k^{j,+}) e_k - \iota_k^{j,+} e_{k+1}$. Moreover $\widehat{G}_j$ has rank N, hence is injective.

Proof. Immediate from (2), whose terms carrying ℓ_k are all proportional to q_k. For injectivity, suppose $\widehat{G}_j c = 0$ and let k be the smallest index with $c_k \neq 0$. Row $k-1$ receives a contribution from column k' only for $k' \in \{k-2, k-1, k\}$, and the first two carry $c_{k'} = 0$. So row $k-1$ reads $-\iota_k^{j,-} c_k = 0$, forcing $c_k = 0$, a contradiction. □

Note that $\widehat{G}_j$ depends only on the grid and on Δ_j, not on the state; the state enters (4) only through the diagonal factor.

Proposition 2.2 (The code-book). *Fix q^0 a probability vector on the grid with mean S_0.*

(i) Given densities q^{j-1} and q^j with $q^j > 0$, the discrete local volatility is uniquely determined by the tridiagonal system

$$G_j(q^j)\, \ell^j \;=\; q^{j-1} - q^j. \tag{5}$$

(ii) Conversely, given $\ell^j > 0$, the vector $q^j := A_j(\ell^j)^{-1} q^{j-1}$ is again a probability vector with mean S_0, and the gridded prices $\pi_i \cdot q^j$ admit no static arbitrage.

(iii) A gridded surface is free of static arbitrage, butterflies and calendar spreads together, if and only if it arises as in (ii) *from some $\ell \geq 0$, and it lies in the* interior *of the arbitrage-free set if and only if $\ell > 0$.*

Proof. (i) is (3), with uniqueness from Lemma 2.1. For (ii), $A_j(\ell^j)^{-1}$ is non-negative with unit column sums and inherits the martingale property [8], attributed there to Andreasen and Huge [1]; non-negativity excludes butterfly arbitrage, and q^j is then a mean-preserving spread of q^{j-1}, so Jensen's inequality excludes calendar arbitrage. (iii) is [8], after [3]: absence of arbitrage is equivalent to non-negative Gamma and Backward-Theta at every node, and ℓ is their ratio. Both conditions are nodewise, and $\ell_i^j = 0$ exactly when one of them vanishes. □

What everything else rests on is that the admissible set is a product of half-lines, one per node, so it has no faces in the sense that a polytope does, and $\lambda = \log \ell$ maps its interior diffeomorphically onto $\mathbb{R}^{NM}$.

2.1 Two conservation laws

The whole structure of the drift restriction follows from one observation about the left null space of G_j. Both identities below are used by [8], where they hold by construction for the backward operator. What is not there, because nothing in that paper needs it, is that $\mathrm{span}\{\mathbf{1}, x\}$ is the *entire* left null space. That is a rank statement, and it is what makes the drift system of Section 3 square.

Lemma 2.3 (Left null space). *For every q and every j,*

$$\mathbf{1}^\top G_j(q) = 0 \qquad \textit{and} \qquad x^\top G_j(q) = 0,$$

where $\mathbf{1} = (1, \ldots, 1)$ and $x = (x_0, \ldots, x_{N+1})$. Consequently $A_j(\ell)$ conserves both total mass and the spot, for any ℓ. Combined with Lemma 2.1, the left null space of $\widehat{G}_j$ is exactly $\mathrm{span}\{\mathbf{1}, x\}$.

Proof. Fix an interior index k and collect the three entries of $G_j(q)$ in column k: from (2) they are $(\iota_k^{j,-}+\iota_k^{j,+})q_k$ in row k, $-\iota_k^{j,+}q_k$ in row $k+1$, and $-\iota_k^{j,-}q_k$ in row $k-1$. Their sum is 0, giving the first identity. Weighting by the node values,

$$q_k\Big[x_k\big(\iota_k^{j,-}+\iota_k^{j,+}\big)-x_{k+1}\iota_k^{j,+}-x_{k-1}\iota_k^{j,-}\Big]=q_k\Big[\iota_k^{j,-}(x_k-x_{k-1})-\iota_k^{j,+}(x_{k+1}-x_k)\Big]=q_k\big[\omega_k^j-\omega_k^j\big]=0$$

by (1), giving the second. The conservation statements follow from (3): $\mathbf{1}^\top A_j(\ell)q=\mathbf{1}^\top q$ and $x^\top A_j(\ell)q=x^\top q$. Since $\widehat{G}_j$ has rank N and lives in $\mathbb{R}^{(N+2)\times N}$, its left null space is two-dimensional, hence exactly $\operatorname{span}\{\mathbf{1},x\}$. □

Remark 2.4 (Counting). $G_j(q)$ has $N+2$ rows and N columns and rank N. Hence the $N+2$ martingale conditions per expiry contain exactly two redundancies, conservation of mass and of the spot, and the remaining N conditions match the N unknowns of ℓ^j. The system is square on the quotient.

3 The drift restriction

Assumption 3.1 (Martingale property)**.** Under $\mathbb{Q}$ every gridded instrument price is a martingale in t. Since the gridded payoffs π_i span $\mathbb{R}^{N+2}$, this is equivalent to each density vector being a martingale,

$$\mathrm{d}q_t^j=\psi_t^j\,\mathrm{d}W_t,\qquad \psi_t^j\in\mathbb{R}^{(N+2)\times d}. \tag{6}$$

This is the only assumption behind the drift restriction, and it carries no regularity. Each q^j is a fixed linear combination of traded butterflies, so (6) is the martingale representation theorem applied to traded prices, and nothing has to commute with a stochastic integral. This is the substantive advantage of a discrete code-book over its continuum analogue, where the corresponding step requires differentiating a stochastic field in strike and maturity and must be licensed by standing regularity conditions. Theorem 3.3 below needs only $q^j>0$.

Theorem 3.2 (Drift restriction)**.** *Let* $\mathrm{d}\ell_t^j=\alpha_t^j\mathrm{d}t+\beta_t^j\mathrm{d}W_t$. *Under Assumption 3.1, for each* $j=1,\dots,M$,

$$\textit{(volatility)}\qquad A_j(\ell^j)\,\psi^j \;=\; \psi^{j-1}-G_j(q^j)\,\beta^j, \tag{7}$$

$$\textit{(drift)}\qquad G_j(q^j)\,\alpha^j \;=\; -\sum_{f=1}^{d}G_j\big(\psi^{j,f}\big)\beta^{j,f}, \tag{8}$$

marched forward in j *from* ψ^0*, the volatility of the initial density, which is determined by the spot dynamics.*

Proof. Apply Itô to $A_j(\ell_t^j)q_t^j=q_t^{j-1}$, using (3) to write the left side as $q^j+G_j(q^j)\ell^j$ and noting that G_j is bilinear. The $\mathrm{d}W$ terms give $\psi^j+G_j(\psi^j)\ell^j+G_j(q^j)\beta^j=\psi^{j-1}$, which is (7) by (3). The $\mathrm{d}t$ terms give $G_j(q^j)\alpha^j+\sum_f G_j(\psi^{j,f})\beta^{j,f}=0$, since q^{j-1} and q^j are driftless by Assumption 3.1 and the only quadratic covariation is between $G_j(\cdot)$ and ℓ^j, both driven by W. □

Equation (7) is not ours. Compared with the pricing recursion $A_j(\ell^j)q^j=q^{j-1}$ it is *the same operator with one added source*, the structure of Carmona and Nadtochiy [9]. The drift restriction below is theirs in the continuum, and Wissel's [19] on this lattice; we claim no priority for either.

What the fixed lattice changes is threefold. Their restriction carries a term in the derivative of the code-book along time to maturity; a lattice of fixed expiry dates has no such derivative, and here that term is absent. The relation is algebraic rather than differential, so neither the high-order regularity their statement requires nor the singular behaviour as time to maturity vanishes that it is built to accommodate can arise here, the front slab having positive width. And the discrete system is over-determined, $N+2$ conditions on N unknowns, which is what the counting of Remark 2.4 squares; that counting is what turns a differential relation into an identity, and it has no continuum counterpart. Appendix A carries both identifications in the original notation.

Theorem 3.3 (The drift is a pointwise covariance)**.** *Define the relative density volatility*

$$\eta_i^j \;:=\; \frac{\psi_i^j}{q_i^j}, \qquad i = 0, \dots, N+1, \tag{9}$$

that is, the relative volatility of q_i^j. Only the interior components enter (10)*, since ℓ lives on the interior nodes; the boundary components are needed in Section 5. Then* (8) *is solvable for every β and every state, its solution is unique, and it is given componentwise by*

$$\alpha_i^j \;=\; -\sum_{f=1}^{d} \eta_i^{j,f}\, \beta_i^{j,f} \qquad \Longleftrightarrow \qquad \mathrm{d}\lambda_i^j = \Big[-\langle \eta_i^j, b_i^j\rangle - \tfrac{1}{2}|b_i^j|^2\Big]\mathrm{d}t + b_i^j\,\mathrm{d}W_t, \tag{10}$$

where $\lambda = \log \ell$ and $b = \beta/\ell$.

Proof. By Lemma 2.1 both sides of (8) factor through the same fixed matrix:

$$\widehat{G}_j \,\mathrm{diag}(q^j)\,\alpha^j \;=\; -\sum_f \widehat{G}_j \,\mathrm{diag}(\psi^{j,f})\,\beta^{j,f}.$$

Since $\widehat{G}_j$ is injective it cancels, leaving $\mathrm{diag}(q^j)\alpha^j = -\sum_f \mathrm{diag}(\psi^{j,f})\beta^{j,f}$, which is (10) after dividing by $q_i^j > 0$. Existence and uniqueness follow from the same cancellation. The logarithmic form is Itô applied to $\lambda = \log \ell$, using $\alpha/\ell = -\langle \eta, b\rangle$. □

Two features of (10) are worth drawing out. The first is its shape. It is the exact analogue of the HJM restriction $\alpha = \sigma \int \sigma$: the drift of the log local variance is minus the covariance of its own volatility with the volatility of the log density, plus an Itô convexity term. Where HJM's derived object is an integral of the volatility across maturities, ours is η, obtained from b by the single forward recursion (7). The second is the cost. Once η is known no matrix is inverted: no Jacobian of the pricing map is ever formed and no rank condition on it appears. And since ψ is linear in β and α is bilinear in (β, ψ), the drift is an exact quadratic form in β, as in HJM; Theorem 5.6 is what rules out the corresponding finite-time explosion.

The driftless point is the classical model: if $\beta \equiv 0$ and $\psi^0 = 0$ then $\psi \equiv 0$ and $\alpha \equiv 0$, so the discrete local variance surface is frozen, and with $\psi^0 \neq 0$ the densities still move, driven by the spot alone, with $\alpha \equiv 0$ persisting.

Proposition 3.4 (Spot consistency)**.** *The spot is $S_t = x^\top q_t^j$, the same for every j by Lemma 2.3. Since $x^\top A_j = x^\top$ and $x^\top \widehat{G}_j = 0$, equation* (7) *gives*

$$\mathbf{1}^\top \psi^j = \mathbf{1}^\top \psi^0 = 0 \qquad \textit{and} \qquad x^\top \psi^j = x^\top \psi^0 = \sigma_t^S S_t \quad \textit{for every } j.$$

So the spot's volatility is the single scalar $x^\top\psi^0$, it is automatically the same at every expiry, and no further condition links the spot to the surface. Taking $\psi^0 = 0$ freezes the spot and recovers $x^\top\psi^j = 0$.

Proposition 3.4 is the discrete form of Carmona and Nadtochiy's spot consistency condition: the spot and its volatility are both *read off* the surface rather than modelled, $S = x^\top q^j$ and $\sigma^S S = x^\top\psi^j$ for any j.

The drift restriction is used below through the following chain, which is what a model actually evaluates:

$$b \xrightarrow{(7)} \psi^1, \dots, \psi^M \xrightarrow{\eta=\psi/q} \alpha, \tag{11}$$

a forward march of d tridiagonal solves per expiry, $O(NMd)$ in total. The volatility structure is the only input; the drift is computed, never fitted.

4 The spot and the front expiry

The spot is not a state variable. By Proposition 3.4 it is the read-off $S_t = x^\top q_t^j$, the same scalar at every expiry. What has to be specified is the front density q^1, the initial condition of the recursion of Proposition 2.2; everything at $j \geq 2$ follows from it and from the volatility structure.

4.1 The front chart

The front density must be strictly positive and sum to one. Both are automatic in the chart

$$q_i^1 = \frac{e^{z_i}}{\sum_k e^{z_k}}, \qquad \mathrm{d}z = m\,\mathrm{d}t + \sum_{f=1}^{d} \nu^f\,\mathrm{d}W^f, \tag{12}$$

with $z \in \mathbb{R}^{N+2}$ unconstrained. The chart is invariant under $z \mapsto z + c\mathbf{1}$, so a shift of ν along $\mathbf{1}$ leaves q^1 unchanged and only the q^1-centred part of ν is identified; write $\mathrm{dev}^f := \nu^f - \langle q^1, \nu^f\rangle$ for it. The input ν is free, and plays the role b plays at $j \geq 2$. Itô's formula for a function of a vector process shows that each q_i^1 is a martingale exactly when

$$m_i = -\tfrac{1}{2}\sum_f \big(\mathrm{dev}_i^f\big)^2 \tag{13}$$

up to a multiple of $\mathbf{1}$. Furthermore we obtain

$$\psi^{1,f} = q^1 \odot \mathrm{dev}^f, \qquad \eta^{1,f} = \mathrm{dev}^f, \qquad \mathbf{1}^\top\psi^{1,f} = 0. \tag{14}$$

The same computation appears in [11] and as the stochastic replicator equation [13].

Remark 4.1 (Resolution at the front expiry). Among grid measures with mean S the least dispersed is the pair of nodes straddling it, so

$$\mathrm{Var}_{q^1}(x) \;\geq\; (S - x_i)(x_{i+1} - S), \qquad x_i \leq S < x_{i+1}, \tag{15}$$

which is at most $\frac{1}{4}h^2$, and the front implied volatility has a floor of order $h/(2S\sqrt{T_1 - t})$. It vanishes when the spot sits on a node, which is exactly what Proposition 4.3 forces at T_1.

4.2 Terminal settlement: pinning the spot to a strike

The chart constrains q^1 at no time $t < T_1$ beyond positivity and mass, so it imposes no terminal condition. Wissel's front parametrisation fixes the terminal state whatever the coefficients [19]; here the terminal state is a modelling choice. We take the front volatility proportional to the distance from the spot, in one factor,

$$\eta^1_{i,t} = \frac{c_t\,(x_i - S_t)}{h}, \qquad c_t = a_1\tau^{-p} \text{ near } T_1, \quad \tau = T_1 - t,\ a_1 > 0. \tag{16}$$

It is already centred, since $\sum_i q^1_i(x_i - S_t) = 0$. Its coefficients are polynomial in q^1 and bounded on $[0, T_1 - \varepsilon]$, so there is a unique strong solution on $[0, T_1)$,

$$q^1_{i,t} = q^1_{i,0}\exp\Big(\int_0^t \eta^1_i\,\mathrm{d}W - \tfrac{1}{2}\int_0^t (\eta^1_i)^2\mathrm{d}s\Big), \qquad \|\eta^1_t\|_\infty \le c_t\,R/h, \quad R := x_{N+1} - x_0. \tag{17}$$

Remark 4.2 (An information-based reading). The structure (16) is the centring of the deterministic volatility $v(t,x) = c_t x/h$. For such structures [11] construct the density process by filtering: draw X from the initial density, observe $I_t = B_t + \int_0^t v(s, X)\,\mathrm{d}s$ with B an independent Brownian motion, and take the conditional law of X. On the lattice, with sums for integrals and by uniqueness in law, $q^1_{i,t}$ is the conditional probability that $X = x_i$, and (17) is their filtering formula written in the innovation W. Pinning is the statement that the observation reveals X at T_1. Their sufficient condition for this requires $\int_0^t c_s\,\mathrm{d}s \to \infty$, which Assumption 5.7 excludes; the window below needs only $\int^{T_1} c_t^2\,\mathrm{d}t = \infty$.

Proposition 4.3 (The front collapses onto a node, and settles exactly)**.** *If $p \ge \frac{1}{2}$ then $q^1_{T_1} := \lim_{t\uparrow T_1} q^1_t = e_k$ for a random index k, almost surely, so $S_{T_1} = x_k$ and*

$$\sum_i q^1_{i,T_1}(x_i - x_n)^+ = (S_{T_1} - x_n)^+ \qquad \textit{for every } n. \tag{18}$$

If $p < 1$ then Assumption 5.7 holds on $[0, T_1]$. The window $p \in [\frac{1}{2}, 1)$ gives both.

Proof. Each q^1_i is a bounded martingale, so $q^1_t \to q^1_{T_1}$ almost surely, and the variance $V_t := \sum_i q^1_{i,t}(x_i - S_t)^2$ converges to that of $q^1_{T_1}$. By (14) and (16) the volatility of $S = x^\top q^1$ is $\sum_i x_i q^1_i \eta^1_i = c_t V_t/h$. Since S is a martingale with values in $[x_0, x_{N+1}]$,

$$\mathbb{E}\int_0^{T_1} \frac{c_t^2 V_t^2}{h^2}\,\mathrm{d}t = \mathbb{E}\big[(S_{T_1} - S_0)^2\big] \le R^2,$$

so the integral is finite almost surely. On the event $\{V_{T_1} > 0\}$ we would have $V_t \ge V_{T_1}/2$ near T_1, and the integral would diverge, because $\int^{T_1} c_t^2\,\mathrm{d}t = a_1^2\int^{T_1}\tau^{-2p}\,\mathrm{d}t = \infty$ for $2p \ge 1$. Hence $V_{T_1} = 0$ almost surely: $q^1_{T_1}$ is a probability vector on distinct points with zero variance, that is e_k with $x_k = S_{T_1}$, and (18) reads $(x_k - x_n)^+ = (S_{T_1} - x_n)^+$. Finally $\int_0^{T_1}\|\eta^1_t\|_\infty \mathrm{d}t \le (R/h)\int_0^{T_1} c_t\,\mathrm{d}t < \infty$ when $p < 1$, by (17). □

The terminal spot is therefore a grid point, as in [19]; the difference is that here it is chosen, by (16), and there it is forced by the parametrisation. Appendix A.2 sets the two charts side by side.

Remark 4.4 (What the lattice could carry, and does not). At expiry each butterfly price q_i^j is worth its payoff: for $S_{T_j} \in [x_k, x_{k+1}]$ the second difference defining q_i^j telescopes to $\varpi(S_{T_j})$, the barycentric weights of the spot on the two nodes straddling it. Since $u \mapsto (u - x_n)^+$ is affine between nodes,

$$\sum_i \varpi_i(s)\,(x_i - x_n)^+ \;=\; (s - x_n)^+ \qquad \text{for every } s \in [x_0, x_{N+1}] \text{ and every model strike } x_n, \tag{19}$$

so $\varpi(s)$ is an admissible terminal state that settles every gridded contract exactly, also for s strictly inside a cell. The chart does not force the pinning; the choice (16) does, because its terminal variance must vanish. A front volatility that drives q^1 to $\varpi(s)$ with s inside a cell is left as future work (Section 7).

5 The wing constant, and existence

The log discrete local variance $\lambda = \log \ell$ ranges over all of $\mathbb{R}^{NM}$, so positivity is automatic. By Theorem 3.3 the model is the stochastic differential equation

$$\mathrm{d}\lambda_t = \Big[-\langle \eta_t, b_t\rangle - \tfrac{1}{2}|b_t|^2\Big]\mathrm{d}t + b_t\,\mathrm{d}W_t, \tag{20}$$

whose only non-elementary coefficient is $\eta = \psi/q$ from (7) and (9). Its denominator is small in the wings, and the question is whether η can blow up there.

The argument follows the route of Wissel's proof of his Theorem 3.5 [19]: log coordinates (his (3.29)), induction over expiries, and an existence result for locally Lipschitz equations (his Proposition 3.7). Its one non-standard ingredient is an a priori bound, independent of the state, on the relative volatility of the calendar spread at the second expiry. Wissel could not obtain this bound in general and imposes his condition (3.24) instead (his Remark 1 following the proof, p. 18); in our coordinates the quantity is $\eta + b$ (Appendix A.2). We obtain it from an exact identity for the volatility recursion and a known inequality for inverses of M-matrices.

5.1 The volatility recursion

Write $P := A_j(\ell^j)^{-1}$, so that $q^j = Pq^{j-1}$.

Corollary 5.1 (Positivity, unit column sums, absorbing ends)**.** *$P \geq 0$ entrywise, $\mathbf{1}^\top P = \mathbf{1}^\top$, so $P_{im} \in [0,1]$, and the outer columns of P are e_0 and e_{N+1}.*

Proof. By (2), A_j has non-positive off-diagonal entries, and in each column the diagonal entry exceeds the sum of the moduli of the off-diagonal ones by exactly one; so A_j is a non-singular M-matrix and $P \geq 0$ [2]. The column sums follow from $\mathbf{1}^\top A_j = \mathbf{1}^\top$ (Lemma 2.3), and the outer columns of A_j are e_0 and e_{N+1}. The first two statements are also Theorem 3.3 of [8], attributed there to [1]. □

That is all that is needed about P itself. What we want next is a recursion for η alone: one step per expiry, each adding a controlled amount, so that the bound telescopes along the ladder. Dividing (7) by q^j gives

$$\eta_i^j \;=\; \frac{(P\psi^{j-1})_i}{q_i^j} \;-\; \frac{\big(P\,G_j(q^j)\beta^j\big)_i}{q_i^j}, \tag{21}$$

the first term carrying η^{j-1} and the second carrying b^j.

The first is an average. Writing $\psi^{j-1} = q^{j-1} \circ \eta^{j-1}$, its weights $P_{ik} q_k^{j-1}/q_i^j$ are non-negative because $P \geq 0$ and sum to one because $q^j = Pq^{j-1}$, so it contributes at most $\|\eta^{j-1}\|_\infty$; they are the matrix $\mathcal{A}$ of (24).

The source is the difficulty. Since G_j is a second difference, row i carries the neighbouring densities $q_{i\pm1}^j$ as well as q_i^j, so dividing by q_i^j leaves the ratios

$$\frac{q_{i-1}^j}{q_i^j} \qquad \text{and} \qquad \frac{q_{i+1}^j}{q_i^j}, \tag{22}$$

which are unbounded in the wings, where the density decays fastest; nothing in (7) bounds them. Let $E \in \mathbb{R}^{N\times(N+2)}$ restrict to the interior nodes, so $EE^\top = I_N$, and recall $A_j = I + \widehat{G}_j \operatorname{diag}(\ell^j) E$.

Proposition 5.2. *With $b^j = \beta^j/\ell^j$ and $\circ$ the entrywise product,*

$$\psi^j \;=\; P\,\psi^{j-1} \;-\; (I-P)\,E^\top\big(q^j \circ b^j\big), \tag{23}$$

equivalently, with the entrywise non-negative

$$\mathcal{A} := \operatorname{diag}(q^j)^{-1} P \operatorname{diag}(q^{j-1}), \qquad \mathcal{B} := \operatorname{diag}(q^j)^{-1} P\, E^\top \operatorname{diag}(Eq^j),$$

$$\eta^j \;=\; \mathcal{A}\,\eta^{j-1} \;+\; \big(\mathcal{B} - E^\top\big) b^j, \tag{24}$$

where $\mathcal{A}\mathbf{1} = \mathbf{1}$, so the transport is an averaging operator, and the i-th row of $\mathcal{B}$ sums to $(Pq^j)_i/q_i^j - \mathbf{1}_{\{i\in\{0,N+1\}\}}$.

Proof. Multiplying $A_j = I + \widehat{G}_j \operatorname{diag}(\ell^j) E$ by P on the left and $E^\top$ on the right gives $P\widehat{G}_j \operatorname{diag}(\ell^j) = (I-P)E^\top$. By (7) and (4), $\psi^j = P\psi^{j-1} - P\widehat{G}_j \operatorname{diag}(q^j)\beta^j$. Since $\beta^j = \ell^j \circ b^j$ and diagonal matrices commute,

$$\operatorname{diag}(q^j)\beta^j = \operatorname{diag}(q^j)\operatorname{diag}(\ell^j)\, b^j = \operatorname{diag}(\ell^j)\operatorname{diag}(q^j)\, b^j = \operatorname{diag}(\ell^j)\big(q^j \circ b^j\big),$$

which is (23). Dividing it by q_i^j gives (24): the rows of $\mathcal{A}$ sum to $(Pq^{j-1})_i/q_i^j = 1$, and for $\mathcal{B}$ the terms $k \in \{0, N+1\}$ missing from $(Pq^j)_i$ vanish for interior i and equal q_i^j at $k = i$ for outer i, by Corollary 5.1. □

So η^j is an average of η^{j-1} plus a source in b^j, and the march across expiries is additive.

Remark 5.3 (Why the source is harmless). The ratios (22) disappear because the source is the operator's own increment: by (3), $G_j(q)\ell = A_j(\ell)q - q$. The mechanism is clearest when $b^j \equiv c$, so that $\beta^j = c\,\ell^j$:

$$G_j(q^j)\beta^j = c\big(A_j q^j - q^j\big) = c\big(q^{j-1} - q^j\big), \qquad P\,G_j(q^j)\beta^j = c\big(q^j - Pq^j\big),$$

by $A_j q^j = q^{j-1}$ and $Pq^{j-1} = q^j$. Divided by q_i^j this is $c\big(1 - (Pq^j)_i/q_i^j\big)$: the two ratios of (22) have collapsed to the single ratio $(Pq^j)_i/q_i^j$, whose size is what Θ_j measures in (25). For general b the factorisation (4) moves b off ℓ and onto q, diagonal matrices commuting, and the same computation applies to $q^j \circ b^j$. This is motivation only: the proof above reaches (23) through $P\widehat{G}_j \operatorname{diag}(\ell^j) = (I-P)E^\top$ in one step.

5.2 The wing constant

That single surviving ratio, $(Pq^j)_i/q_i^j$, is what has to be bounded. Write

$$R_i^j := \mathbf{1}_{\{1\le i\le N\}} + \frac{(Pq^j)_i}{q_i^j} - \mathbf{1}_{\{i\in\{0,N+1\}\}}, \qquad \Theta_j := \max_i R_i^j, \qquad \Xi := \sum_{r=1}^{N} P_{rr} \le N, \tag{25}$$

and $H_j := \max_i|\eta_i^j|$. Here R_i^j is the row bound of (24) at node i, and Θ_j, the *wing constant*, is its size across the slab. The theorem shows Θ_j is bounded by the grid alone, uniformly in the state; the bound rests on one inequality for inverses of M-matrices and its consequence for P^2.

Lemma 5.4 (Path product inequality)**.** *For all nodes i, r, m,*

$$P_{ir}\,P_{rm} \;\le\; P_{im}\,P_{rr}. \tag{26}$$

Proof. P is the inverse of an M-matrix by Corollary 5.1, and every inverse M-matrix satisfies the path product conditions $a_{ij}a_{jk} \le a_{ik}a_{jj}$ for all i, j, k [20], a result first noted by Willoughby [18] and proved by Johnson and Smith [12]. Take $(i, j, k) = (i, r, m)$. □

Corollary 5.5 (Squared propagator)**.** *For all i, m,*

$$\big(P^2\big)_{im} \le \begin{cases} \Xi\,P_{im}, & 1 \le i \le N, \\ (1+\Xi)\,P_{im}, & i \in \{0, N+1\}. \end{cases} \tag{27}$$

Proof. Sum (26) over r. An outer r contributes only if $r = i$, since column r of P is e_r, and then contributes P_{im}. □

Theorem 5.6. *For every $\ell \in (0,\infty)^N$,*

$$H_j \;\le\; H_{j-1} + \Theta_j\|b^j\|_\infty, \tag{28}$$

and

$$\Theta_j \le 1 + \Xi \le N + 1, \qquad \Theta^* := \sup_\ell \sum_{j=2}^{M} \Theta_j \le (M-1)(N+1). \tag{29}$$

Proof. Since $q^j = Pq^{j-1} \ge 0$, Corollary 5.5 gives $(Pq^j)_i = \sum_m (P^2)_{im} q_m^{j-1} \le \Xi\, q_i^j$ for interior i and $\le (1+\Xi) q_i^j$ for outer i, so $R_i^j \le 1 + \Xi$ in both cases. Finally, $\mathcal{A} \ge 0$ with $\mathcal{A}\mathbf{1} = \mathbf{1}$ gives $\|\mathcal{A}\eta^{j-1}\|_\infty \le H_{j-1}$, while row i of $(\mathcal{B} - E^\top)b^j$ has modulus at most $\|b^j\|_\infty R_i^j$ by the row sums of Proposition 5.2. □

The bound is a property of the grid, not of ℓ, and it is far from tight on realistic states: on Heston surfaces with $N = 40$ the measured Θ_j is at most 7.6 against 41.

5.3 Existence

Iterating (28) from the front slice,

$$\sup_{j,i}|\eta_{i,t}^j| \;\le\; \|\eta_t^1\|_\infty + (M-1)(N+1)\,\|b\|_\infty, \tag{30}$$

and iterating (23) shows that η^j is affine in η^1, with coefficients depending on λ only through $P_2, \dots, P_j$ and $q^2, \dots, q^j$. The only t on the right of (30) is inside η_t^1, which enters linearly, so it has to be integrable in time rather than bounded.

Assumption 5.7 (Integrable front volatility)**.** For a horizon $T \le T_1$,

$$\int_0^T \|\eta_t^1\|_\infty \, \mathrm{d}t \; < \; \infty \qquad \text{almost surely.} \tag{31}$$

TBD.

6 Finite-step simulation and numerical evidence

The drift restriction of Section 3 admits two numerically distinct implementations. The first integrates the log-DLV SDE directly. It is the closest implementation of the continuous-time model, but a finite time step is only approximately martingale preserving. The second advances the probability density itself through a finite set of symmetric branches. It is an exact one-step martingale construction on the modeled finite support, although it is not a claim about a continuous-strike or unbounded-support limit. We report both because their numerical guarantees differ.

All experiments in this section use four weekly expiries and fixed cash strikes within each weekly transition. "DLVsq" denotes the positive discrete local variance ℓ. The final density-branch audit uses the discrete factor X alone. Earlier development and long-horizon audits also multiply X by an independent mean-one Black factor Y with annual volatility 10%; the conditioning identities below apply to both, using the Black semigroup for Y.

6.1 Direct integration of the DLV diffusion

For the direct SDE experiment we work in log excess-DLV coordinates

$$z^{j,i} = \log(\ell^{j,i} - \ell_{\min}).$$

The volatility of z is the rank-25 residual covariance estimated from 703 weekly historical transitions through 31 December 2024. At every substep the density volatility is obtained from (7), and the drift of ℓ is then computed from (10), including the Itô term in log coordinates. Each week is integrated by ten stochastic-Heun substeps, using the same Gaussian increment in predictor and corrector. At the weekly endpoint spot is drawn from the front density and the expired slab is removed.

The audit starts from historical state 44 (8 March 2022), uses 200 interior strikes, 100,000 paths arranged as 50,000 antithetic pairs, and follows a shrinking tower for four weeks without re-gridding. Fifty antithetic pairs are aggregated into each of 1,000 independent randomized-stratification blocks. There are 800, 600, 400, and 200 live fixed-contract option returns at the four trade dates, hence 2,000 coordinates in the joint test.

For this experiment the blockwise coordinate statistics are supplemented by a 10,000-draw shared-sign multiplier bootstrap. If t_k is the usual block t-statistic, the two simultaneous statistics are

$$T_{\max} = \max_k |t_k|, \qquad T_{\mathrm{RMS}} = \left(d^{-1} \sum_{k=1}^{d} t_k^2 \right)^{1/2}.$$

The reported decision value is $p_{\mathrm{comb}} = \min\{1, 2\min(p_{\max}, p_{\mathrm{RMS}})\}$. The shared multiplier for all coordinates of a block preserves their observed cross-option dependence.

Audit	Active	$\max \lvert t \rvert$	RMS t	p_{comb}	5% decision
One-step option returns	2,000	3.334	1.109	0.2746	do not reject
Terminal option returns	2,000	3.334	1.044	0.3210	do not reject
Terminal option price changes	2,000	2.934	1.156	0.3500	do not reject
Integrated-DLV control	1,200	1.950	0.878	1.0000	do not reject
Weekly spot returns	4	2.431	1.561	0.0876	do not reject

Table 1: Martingale diagnostics for the ten-substep Heun implementation. The 800 omitted integrated-DLV coordinates have effectively zero block variance. The primary one-step audit has separate multiplier-bootstrap values $p_{\max} = 0.1373$ and $p_{\text{RMS}} = 0.2465$; its simultaneous 5% max-t cutoff is 4.235.

The largest density-mass residual was 9.33×10^{-14} and the largest conditional-forward residual was 8.93×10^{-14}. The minimum density was zero, due to finite-support point masses rather than a negative density. The simulation took 72.1 seconds, or approximately 1,388 paths per second. These results support the discretization, but they do not turn an infinitesimal drift restriction into an exact one-week transition.

6.2 An exact finite-step density construction

We next construct the one-week law in density coordinates. Let $\varepsilon_b \in \{-1, +1\}^d$ be the b-th Rademacher branch and let $\pi_b = 2^{-d}$, with $d = 4$ in the experiment. Write Γ^j for the state-dependent loading of density q^j implied by the chosen DLVsq loading. The branch marginal is

$$q^{j,b} = q^j + \sqrt{h}\,\Gamma^j \varepsilon_b, \qquad b = 1, \ldots, 2^d. \tag{32}$$

Since $\sum_b \pi_b \varepsilon_b = 0$ and $\sum_b \pi_b \varepsilon_b \varepsilon_b^\top = I_d$,

$$\sum_b \pi_b q^{j,b} = q^j \tag{33}$$

holds exactly, before Monte Carlo sampling. DLVsq is then recovered algebraically from each $q^{j,b}$ by Proposition 2.2; it is not advanced by an Euler step or fitted by a nonlinear optimizer. A branch is accepted only if its masses are non-negative and the recovered DLVsq lies in its configured open bounds. The implementation checks mass, mean, decoder consistency, and DLV bounds at every step.

In the implemented weekly law the front density is branch invariant, $q^{1,b} = q^1$. Independently of b, a front node $I = i$ is drawn with probability q_i^1. For every surviving expiry $j \geq 2$, let $R_{ir}^{b,j}$ be the conditional probability of node r given branch b and front node i. The SANOS transition supplies rows satisfying

$$\sum_r R_{ir}^{b,j} = 1, \qquad \sum_r x_r R_{ir}^{b,j} = x_i, \qquad \sum_i q_i^1 R_{ir}^{b,j} = q_r^{j,b}. \tag{34}$$

The first two identities make every row a martingale transition; the third recovers the branch marginal. Together with (33), they imply the tower property for every payoff represented on the finite support.

This construction is exact on the finite model grid. It is not asserted to be the unique risk-neutral law, and the particular independent sampling of b and I does not generate contemporaneous correlation between the raw surface branch and front-node innovation. Realized

density-weighted factor covariance nevertheless matched the intended empirical covariance closely: in the independent all-state confirmation the largest relative error was 1.87% and the RMS error was 0.55%.

6.3 Conditioning and re-gridding

Suppose current cash spot is s and the revealed front node is x_i, so the next X spot is sx_i. Define the next relative support by

$$x_r'^{(i)} = \frac{x_r}{x_i}. \tag{35}$$

Then $(sx_i)x_r'^{(i)} = sx_r$ path by path. Consequently the price of a surviving fixed-cash call is identical on the inherited and re-gridded coordinates,

$$C_{t+h}^{b,i,j}(K) = \sum_r R_{ir}^{b,j}(sx_r - K)^+.$$

Averaging first over i and then over b gives

$$\begin{aligned}\mathbb{E}_t[C_{t+h}^j(K)] &= \sum_b \pi_b \sum_i q_i^1 \sum_r R_{ir}^{b,j}(sx_r - K)^+ \\ &= \sum_r q_r^j (sx_r - K)^+ = C_t^j(K).\end{aligned}$$

Thus re-gridding is an exact coordinate change for all surviving fixed-cash claims, not an interpolation. With an independent Black factor Y, the same proof conditions on Y and then applies its mean-one Black semigroup. A newly issued far expiry is different: there is no pre-existing fifth-week claim whose price it must preserve.

At each weekly roll, the old 2W–4W densities become the new floating 1W–3W densities through (34) and (35). Only the new 4W DLVsq curve is supplied statistically. The retained boundary model is a 25-factor PCA followed by an unrestricted weekly VAR(1) in factor scores. It uses 703 training pairs through 2024, explains 99.9912% of the 4W curve variation, and has spectral radius 0.9089. On 155 untouched 2025 pairs its raw-coordinate RMSE is 0.5902, against 0.6526 for persistence, a 9.56% improvement. SANOS constructs the new 4W density as a static martingale extension of the inherited 3W density.

6.4 Martingale audit for the exact density step

Adjacent simulated paths form an antithetic pair. For pair $a = 1, \dots, n$ and audited return coordinate k, let R_{ak} be the pair mean, $\bar{R}$ the vector of sample means, and $\widehat{\Sigma}$ their sample covariance. We report

$$z_k = \frac{\sqrt{n}\,\bar{R}_k}{\widehat{\sigma}_k}, \qquad J = n\bar{R}^\top \widehat{\Sigma}^+ \bar{R}. \tag{36}$$

The pseudo-inverse is formed in correlation space after discarding eigenmodes at or below 0.01. If ν modes remain, the reference value is $J \sim \chi^2_\nu$ under the joint zero-mean null. A small upper-tail p-value rejects zero drift; a large value means only that the observed drift is compatible with Monte Carlo noise. We use 0.05 for one look, 0.0125 for four weekly looks, and 0.003125 for 16 looks. Because the options are highly collinear and the retained modes are estimated from the sample, the exact density-tower identities and matched cash/re-grid prices are stronger structural diagnostics than an isolated asymptotic p-value.

The canonical X-only audit used all 891 certified historical states, 100,000 paths (50,000 independent antithetic pairs), 11 ITM call/put strikes per expiry, four one-week steps, 64-bit arithmetic, and 10% annualized log-DLVsq volatility. Each state received 56 or 57 independent pairs. The audited book contains spot and the surviving original fixed-cash claims.

Week	Worst $\lvert z\rvert$	Max option mean (bp)	J	ν	p
1	1.796	0.497	10.980	7	0.139
2	2.143	0.863	7.228	7	0.406
3	1.371	0.696	5.858	6	0.439
4	2.057	1.245	5.506	5	0.357

Table 2: All-state X-only martingale audit for the surviving original fixed-cash book plus spot. Option means are basis points of initial spot. No horizon rejects the joint zero-mean null at 5%.

The deterministic checks close much more tightly than Monte Carlo sampling: the maximum density-tower residual is 6.49×10^{-15}, option-tower residual 1.67×10^{-16}, spot-mean residual 2.60×10^{-14}, four-roll decoder residual 1.77×10^{-13}, and first re-grid cash-price residual 3.33×10^{-16}.

A matched run decomposed every weekly change into the fixed-expiry surface branch and the spot-conditioning/re-grid component, using identical random draws. Direct inherited-cash and re-gridded valuations agree pathwise to at most 2.22×10^{-16} and have identical $z/J/p$ statistics.

Week	Fixed-expiry p	Full first-member p	Full pair p	Max price difference
1	0.366	0.684	0.139	0
2	0.538	0.163	0.128	2.22×10^{-16}
3	0.911	0.414	0.0208	2.22×10^{-16}
4	0.945	0.0808	0.546	2.22×10^{-16}

Table 3: Matched fixed-support and floating-coordinate decomposition. The first two p-value columns use one independent member of each pair; the third uses antithetic pair means.

The week-3 pair-mean value is a nominal rejection at 5%, but not at the pre-specified four-look cutoff 0.0125. It is also visibly sensitive to the correlation eigenvalue cutoff: its p-value is 0.0429, 0.0208, 0.0879, and 0.295 at cutoffs 0.005, 0.01, 0.02, and 0.05. Estimating covariance and the mean on opposite half-samples gives 0.214 and 0.338. No coordinate has $\lvert z\rvert > 2$ and the largest mean is 0.363bp. Since the direct and re-gridded path matrices are identical, this fluctuation cannot be caused by re-gridding. The fixed-expiry branch alone has maximum option means of only 0.0046, 0.0056, 0.0013, and 0.0050bp over weeks 1–4.

Earlier development audits reached the same conclusion under a nonzero Black backbone. On 21 states with 10,000 paths per state, the four weekly maximum coordinate $\lvert z\rvert$ values were 2.400, 3.203, 2.871, and 2.968, while the covariance-aware Fisher-combined state-test p-values were 0.871, 0.672, 0.852, and 0.522. A heterogeneous 10,000-path confirmation covering all 891 states gave unconditional maximum $\lvert z\rvert$ values 1.575, 1.970, 1.878, and 1.040 and p-values 0.702, 0.414, 0.067, and 0.317. Its high-dimensional state-feature-weighted values were 0.027, 0.759, 0.251, and 0.0089; they are reported because they expose the poor calibration of hard spectral cutoffs in small per-state samples, rather than being hidden behind the aggregate result.

6.5 Repeated rolls, speed, and finite-support behavior

A 16-week audit with the independent 10% Black backbone used 10,000 paths and all 891 states. Across all rolls the largest decoder residual was 2.35×10^{-13}, conditional mass and mean errors stayed below 6.6×10^{-14}, and DLVsq remained strictly inside $(10^{-4}, 9)$. The maximum spot $|z|$ over 16 horizons was 2.126, the maximum rolling-option $|z|$ was 2.417, and every rolling-option joint p-value was at least 0.0839. With the last-expiry AR boundary, median cross-expiry dispersion of $\sqrt{\mathrm{DLVsq}}$ was 0.0459 after week 1 and 0.0569 after week 16, so the term structure did not collapse. The corresponding implementation ran 160,000 path-week steps in 14.6 seconds.

The exact branch method is materially faster than a nonlinear finite-step solve. In the heterogeneous 10,000-path benchmark it processed 40,000 path-week steps in 8.84 seconds (4,527 per second), including 7.36 seconds to construct 5,000 distinct initial kernels. This was about 23 times faster than the earlier finite-h Newton transition (195 per second), but about 5.7 times slower than the local SDE transition (25,649 per second). The final X-only audit, with more extensive all-state diagnostics, processed 400,000 path-week steps in 169.8 seconds. These timings are indicative rather than a controlled hardware benchmark.

The remaining limitation is economic rather than a floating-point failure. If the lowest support node x_0 is revealed, then

$$\sum_r R^{b,j}_{0r} x_r = x_0, \qquad x_r \geq x_0,$$

forces $R^{b,j}_{00} = 1$. The upper endpoint is analogous. The X law is then absorbing, all surviving options are intrinsic, and interior DLVsq coordinates are price-unidentified. The positive-density formula $\eta = \psi/q$ no longer applies at zero-mass nodes; any continued DLV coordinates there are an economically invisible numerical convention.

In the four-week X-only run, boundary incidence was 0.003%, 0.003%, 0.003%, and 0.006%. In the 16-week run, the fully collapsed fraction grew from 0.01% to 0.48%; the median week-16 front density retained 23.44 effective nodes, while its first percentile retained 1.15. Avoiding this absorbing tail requires dynamic support expansion, an explicit tail continuation, or a boundary state carrying a conditional tail. Re-gridding alone cannot create support beyond an extreme node while preserving the inherited option surface.

The numerical evidence therefore supports three distinct statements: the Heun implementation is a credible approximation to the continuous density drift; the symmetric density construction is exactly martingale consistent on the finite support; and weekly re-gridding preserves every surviving fixed-cash price path by path. It does not establish uniqueness, an unbounded-support model, or the absence of boundary bias at horizons much longer than those tested.

7 What is open

1. *An off-lattice terminal spot.* By Remark 4.4 the chart represents every $s \in [x_0, x_{N+1}]$ at expiry, and the state that does so is $\varpi(s)$, carried by the two nodes straddling s. Whether a front volatility exists driving q^1 to $\varpi(s)$ for s strictly inside a cell, rather than to a single node as (16) does, is open. The obstruction is that exact settlement forces $q^1_{T_1} = \varpi(S_{T_1})$ and hence $\int_0^{T_1} (\eta^1_i)^2 \mathrm{d}t < \infty$ for both surviving nodes, which by the argument of Proposition 4.3 they cannot both satisfy when the volatility is proportional to distance from the spot. Until it is answered the model of Proposition 4.3 pins, as [19] does, and the question he raises remains open.

2. *The continuum limit.* The constant $(M-1)(N+1)$ is sharp as an inequality and grows with the grid. Whether $\sup_N \Theta_j < \infty$ on a class of realistic surfaces, which is what a continuum statement would need, is open.

References

[1] Jesper Andreasen and Brian Norsk Huge. Volatility interpolation. Available at SSRN 1694972, 2010. Also published in Risk, March 2011.

[2] Abraham Berman and Robert J. Plemmons. *Nonnegative Matrices in the Mathematical Sciences*, volume 9 of *Classics in Applied Mathematics*. SIAM, 1994. Reprint of the 1979 Academic Press edition.

[3] Hans Buehler. Expensive martingales. *Quantitative Finance*, 6(3):207–218, 2006.

[4] Hans Buehler. Volatility and dividends - volatility modelling with cash dividends and simple credit risk. `https://papers.ssrn.com/abstract=1141877`, 2008.

[5] Hans Buehler, Lukas Gonon, Josef Teichmann, and Ben Wood. Deep hedging. *Quantitative Finance*, 19:1271–1291, 2019. SSRN working paper version: doi:10.2139/ssrn.3120710 (2018).

[6] Hans Buehler, Blanka Horvath, Anastais Kratsios, Yannick Limmer, and Raeid Saqur. Sanos - smooth arbitrage-free non-parametric option surfaces. `https://arxiv.org/abs/2601.11209`, Februar 2026.

[7] Hans Buehler, Blanka Horvath, and Anastasis Kratsios. Dysanos generative dynamic smooth arbitrage-free non-parametric option surfaces, 2026.

[8] Hans Buehler and Evgenii Ryskin. Discrete local volatility for large time steps (short version). Technical report, SSRN, 2016. Short version. An extended version also exists: SSRN 2642630 (2015), doi:10.2139/ssrn.2642630.

[9] René Carmona and Sergey Nadtochiy. Local volatility dynamic models. *Finance and Stochastics*, 13:1–48, 2009.

[10] Samuel N. Cohen, Christoph Reisinger, and Sheng Wang. Arbitrage-free neural-sde market models, 2021. arXiv:2105.11053.

[11] Damir Filipović, Lane P. Hughston, and Andrea Macrina. Conditional density models for asset pricing. *International Journal of Theoretical and Applied Finance*, 15(1):1250002, 2012.

[12] Charles R. Johnson and Ronald L. Smith. Path product matrices. *Linear and Multilinear Algebra*, 46:177–191, 1999.

[13] Panayotis Mertikopoulos and Aris L. Moustakas. The emergence of rational behavior in the presence of stochastic perturbations. *The Annals of Applied Probability*, 20(4):1359–1388, 2010.

[14] Martin Schweizer and Johannes Wissel. Arbitrage-free market models for option prices: The multi-strike case. *Finance and Stochastics*, 12:469–505, 2008.

[15] Martin Schweizer and Johannes Wissel. Term structures of implied volatilities: Absence of arbitrage and existence results. *Mathematical Finance*, 18:77–114, 2008.

[16] Magnus Wiese, Lianjun Bai, Ben Wood, and Hans Buehler. Deep hedging: Learning to simulate equity option markets, 2019. NeurIPS 2019 Workshop on Robust AI in Financial Services.

[17] Magnus Wiese, Ben Wood, Alexandre Pachoud, Ralf Korn, Hans Buehler, Phillip Murray, and Lianjun Bai. Multi-asset spot and option market simulation, 2021. arXiv:2112.06823.

[18] R. A. Willoughby. The inverse M-matrix problem. *Linear Algebra and its Applications*, 18:75–94, 1977.

[19] Johannes Wissel. Arbitrage-free market models for option prices. NCCR FINRISK Working Paper 428, NCCR FINRISK, ETH Zürich, 2007.

[20] Yan Zhu, Cheng-Yi Zhang, and Jun Liu. Path product and inverse M-matrices. *Electronic Journal of Linear Algebra*, 22:644–652, 2011.

A Relation to the continuum and lattice results

TBD.

A.1 Carmona and Nadtochiy

TBD.

A.2 Wissel

TBD.

Remark A.1 (The quantile chart forces the spot onto the grid). *TBD.*