# Stochastic Discrete Local Volatility

Magnus Wiese  Hans Buehler*  Blanka Horvath*

*Mathematical Institute and Oxford-Man Institute, University of Oxford

October 3, 2026

**Abstract**

We present a fully arbitrage-free option market model using the SANOS [6] framework based in ideas of [19].

The model is the product of a risk-neutral discrete state martingale defined on a fixed grid with stochastic "discrete local volatilities", and a simple SANOS-type Black & Scholes smoothing process which allows smooth interpolation of option prices between grid points.

We train the model on S&P 500 Index options data.

**Abstract**

We construct an option market model on a fixed lattice of cash strikes and expiries that is free of static and dynamic arbitrage by construction. The state is the vector of gridded risk-neutral densities, encoded by discrete local variances: their positivity is equivalent to the absence of static arbitrage, so in log coordinates the state space is Euclidean. Absence of dynamic arbitrage then fixes the drift node by node, as a covariance between the volatility of each discrete local variance and the relative volatility of the density at the same node. The restriction goes back to Carmona and Nadtochiy [9] in the continuum and to Wissel [19] on this lattice; on a fixed lattice it becomes an algebraic identity, free of the regularity conditions and the maturity derivative of the continuum, and is computed by one tridiagonal solve per factor and expiry. The spot is read off the densities, and at each expiry the front density can be made to collapse onto a node, so every gridded contract settles at its exact payoff. Our main result is existence: for any bounded, locally Lipschitz volatility structure and any number of factors, the model has a unique strong solution up to the first expiry and, restarting at each expiry, across all expiries, with no conditions on the wings. The key estimate bounds the growth of the relative density volatility $\eta = \psi/q$ across expiries by a constant of the grid alone, turning an assumption of Wissel into a theorem.

## 1 Introduction

This article presents an arbitrage-free market model for for smooth surfaces of European option prices.

The core architecture follows the SANOS approach [6] which models the stock price as the product of a time-and-state discrete process $X$ and a smoothing process $Y$, in our case Black & Scholes. The former is used to encode the option market using a discrete form of local volatilities as code book. That does not mean that the model is a local volatility model – we merely use the (discrete) local volatilities as a convenient representation of the option surface. Compared to using implied volatilities the advantage is that mere positivity of the local volatilities is sufficient to ensure the the option market is free of static arbitrage at all times. We have shown

in [6] that this representation can be fitted to real life market data very efficiently using linear programming.

In [7] we have shown how to estimate the dynamics such (discrete) local volatilities of such a model under the statistical measure $\mathbb{P}$ which fit very well to real-life observed option bid/asks.

The key contribution in this paper is showing how to construct a corresponding risk-neutral measure $\mathbb{Q}$ or, equivalently, the risk-neutral drift of the model. Our approach is closely related to the seminal but unpublished work [19].

An option market model reverses the usual direction of construction: instead of specifying a process for the underlying and deriving option prices from it, one models the code-book of prices directly, so that a calibration is a change of coordinates rather than an inverse problem. Two conditions must then be maintained at once. The state must remain free of *static* arbitrage, and its dynamics must admit one equivalent local martingale measure common to the underlying and to every traded option. The second is the option analogue of an HJM drift restriction: the volatility structure may be specified freely, and the drift is then determined. The lattice here is *finite*, finitely many strikes and finitely many expiry dates, which is what is traded.

**What is already known.** Schweizer and Wissel treat the problem in the continuum in two companion papers, one in each direction of the surface: [16] varies maturity at a single strike, and [15] varies strike at a single maturity. In each the code-book is chosen so that its positivity encodes absence of arbitrage, the drift restriction is derived, and existence is proved. The multi-strike multi-maturity case, which is ours, is stated as open in [15], with the expectation that a new parametrisation will be needed. Carmona and Nadtochiy [9] give a continuum local-volatility code-book with a drift restriction, but no existence theorem. On a finite lattice, and answering the multi-maturity question, stands Wissel [19]; his construction is the one this paper is closest to, and it is described separately below.

**Modern approaches.** The question was reopened from the data side, by the scarcity of option-market histories for training and evaluating hedging strategies [5]. Cohen, Reisinger and Wang [10] build neural stochastic differential equations on a floating lattice and penalise a least-squares residual in place of the drift restriction. [17] simulate an equity option book with a generative adversarial network whose state is the discrete local variance grid compressed by principal components, so that static arbitrage is absent by positivity of that grid; that is the code-book choice this paper inherits, and [18] extend it to a multi-asset universe. [7] decode statically arbitrage-free surfaces from a low-dimensional code, fit a linear state process to it, and test for dynamic arbitrage after the fact, reporting it in 12.5% of states at zero transaction cost.

None of the four enforces the drift restriction exactly, and that is where this paper starts: the cross-sectional constraints govern each date's surface and its terminal payoff cone, and leave the transition of a fixed contract between two dates unconstrained. The choice of grid is part of the same point. All four index their nodes relative to the spot rather than in fixed cash terms, and on such a grid a node is not one contract through time, so there is no single instrument whose price is required to be a martingale and nothing for an HJM-type restriction to constrain. Fixing the strikes in cash is what buys the restriction here.

**The closest prior work.** Wissel [19] works on a finite lattice of fixed cash strikes and fixed expiry dates, parametrises it by a price level together with positive local implied volatilities, derives the no-arbitrage drift restrictions, and proves constructive strong existence and uniqueness

for a class of coefficients. Behind the front expiry his local implied variance is *exactly* the backward discrete local variance used here, an algebraic identity on the grid, and his drift restriction is, after a change of measure, exactly the one we state in Theorem 3.2. We claim no priority for it. That working paper appears never to have been published in archival form.

## What this paper does

We keep the same lattice and change the chart: the state is the vector of gridded *densities* together with the discrete local variances, rather than prices together with quantiles of digitals. Three consequences follow, and they are the contributions.

1. **The spot is a read-off, and it settles exactly.** The spot is $S_t = K^\top q_t^j$ at every expiry $j$ (Proposition 3.3), an identity of the chart holding at all $t$. At the front expiry Proposition 4.2 exhibits a family of front volatilities under which $q^1$ collapses onto a single node, so the spot lands on a strike and every gridded contract settles at its exact payoff (Section 4.2). That pinning is a property of the volatility chosen and not of the chart: by Remark 4.3 the terminal states the chart admits carry every spot in the grid range, whereas in [19] the pinning follows from the chart whatever the coefficients. We do not construct dynamics reaching an off-lattice terminal spot; that is left as future work.

2. **Existence, with no condition on the wings.** For every bounded and locally Lipschitz volatility structure $b$ and any number of driving factors, the model has a unique strong solution on $[0, T]$ for every $T < T_1$, positive by construction; it extends to $T_1$ (Theorem 5.6(ii)) and, restarting at each expiry with the settling front, to the whole ladder $[0, T_M]$, unique given that front (Corollary 5.8). The quantity that must be controlled is $\eta = \psi/q$, the density volatility relative to the density, and its growth across expiries is bounded by the grid alone: Theorem 5.4 gives $\Theta_j \le N + 1$ uniformly in the state and in the volatility structure. *Wissel [19] must assume a bound of this kind; here it is a theorem.* The front expiry costs one hypothesis, and it is a mild one: $\eta^1$ must be integrable in $t$, not bounded, because $\eta^j$ is affine in $\eta^1$ and an Itô equation needs an integrable drift rather than a bounded one. The front volatility used here diverges at rate $(T_1 - t)^{-p}$, and the two demands made of $p$ leave a window: $p < 1$ is what Assumption 5.5 asks, while $p \ge \frac{1}{2}$ is what makes the front density collapse, so that every gridded contract settles at its exact payoff (Proposition 4.2).

3. **The drift is pointwise, and the count works out.** In density coordinates the drift equation is $N + 2$ conditions on $N$ unknowns, and is consistent because the left null space of $\widehat{G}_j$ is exactly $\operatorname{span}\{\mathbf{1}, K\}$, conservation of mass and of the spot (Lemma 2.4). The state-independent factor then cancels from both sides and the drift is available nodewise, $\alpha^{j,i} = -\langle \eta^{j,i}, \beta^{j,i} \rangle$, with no matrix inversion and no rank condition.

## Scope, and what is not claimed

The over-determination of the third contribution is not intrinsic: it is created by the choice of density coordinates and removed by the same choice, and in price coordinates the system is square and the question does not arise.

The existence statement is for a fixed finite grid. The constant is $(M-1)(N+1)$ and depends on the grid. The one hypothesis, Assumption 5.5, is a condition on the front volatility and not on the state, and it is sufficient rather than necessary.

## 2 The code-book

We assume deterministic forwards and discount factors, which reduces to a driftless spot with unit discount factors, cf. [4]. The spot at $t = 0$ is $S_0$.

The approach we are presenting relies on a combination of a time-and-state backbone process $X$ and an additional smoothing process $Y$, following the [6] approach to smoothing. However, contrary to that model here we present a full stochastic local volatility model. The weight of the model is carried by the former and echos ideas of [19].

We will define the backbone process on a grid of fixed expiries $0 = T_0 < T_1 < \cdots < T_M$ and strikes $0 < K^0 < K^1 < \cdots < S_0 < \cdots < K^N < K^{N+1}$. We will use two time indices: $t$ will denote model time and $j$ will index fixed expires. We will use $i$ to index strikes Contrary to [6] and similar to [19] expiries and strikes are not floating; they are fixed. The model will evolve along the grid. Specifically, at time $t$ the future expiries in the model are those $T_j$ such that $T_j > t$. We assume that the largest strike spacing is $h := \max_i(K^{i+1} - K^i)$.

### 2.1 Discrete Local Volatility Backbone

At time $t$ we denote by $\Sigma_t^j = (\Sigma_t^{j,i})_{i=1,\dots,N}$ the squared discrete local volatility at the interior nodes $i = 1, \dots, N$ for expiry $T_j > t$. We let $\lambda = \log \Sigma$ throughout. Define

$$\omega^k := \frac{(K^k)^2}{K^{k+1} - K^{k-1}}, \qquad \iota^{k,-} := \frac{\omega^k}{K^k - K^{k-1}}, \qquad \iota^{k,+} := \frac{\omega^k}{K^{k+1} - K^k}, \qquad k = 1, \dots, N, \quad (1)$$

and then the tridiagonal matrix $A_j(\Sigma)$ with rows

$$\big(A_j(\Sigma)q\big)_i = \begin{cases} q^0 - \Delta_j \Sigma^1 \iota^{1,-} q^1, & i = 0, \\ -\Delta_j \Sigma^{i-1} \iota^{i-1,+} q^{i-1} + \big(1 + \Delta_j \Sigma^i (\iota^{i,-} + \iota^{i,+})\big) q^i - \Delta_j \Sigma^{i+1} \iota^{i+1,-} q^{i+1}, & 1 \le i \le N, \\ -\Delta_j \Sigma^N \iota^{N,+} q^N + q^{N+1}, & i = N+1. \end{cases} \quad (2)$$

(here $q^i$ denotes the $i$-th component of a generic $q \in \mathbb{R}^{N+2}$). Up to notation this is the implicit forward operator of [8].

Let now $X_t$ be the current spot price at time $t$ for the backbone process $X$ and let $q_t^0 \in \mathbb{R}^{N+2}$ the density

$$q_t^{0,i} = \frac{(X_t - K^{i-1})^+ - (X_t - K^i)^+}{K^i - K^{i-1}} - \frac{(X_t - K^i)^+ - (X_t - K^{i+1})^+}{K^{i+1} - K^i}$$

Then, the family $q_t = (q_t^j)_{j=0,\dots,M}$ starting in $q_t^0$ is then given iteratively as

$$A_j(\Sigma_t^j)\, q_t^j = q_t^{j-1} \quad (3)$$

Note that in (2) the entry multiplying $q_t^{i\pm1}$ carries $\Sigma_t^{i\pm1}$ and not $\Sigma_t^i$. The following was shown in [8]:

**Theorem 2.1.** *The family is then $q = (q_t^{j,i})$ with $q_t^j \in \mathbb{R}^{N+2}$ us a discrete risk-neutral density on the grid $(T, K)$ in the sense that*

1. $q_t^{j,i} \ge 0$
2. $\sum_{i=0}^{N+1} q_t^{j,i} = 1$

3. $\sum_{i=0}^{N+1} K^i q_t^{j,i} = X_0$.

*This means there is a discrete time-and-space process $X_t = (X_t^j)_{j=0,\dots,M}$ which takes values only in $(T, K)$. Note that the model-time index here is $j$. Call prices for $X_t$ for arbitrary $(\tau, k) \in \mathbb{R}^2_{\geq 0}$ are given by the linear functional*

$$C_t(k,\tau) = \mathbb{E}\left[\left(X_t^j - k\right)^+\right] = \sum_{i=0}^{N+1} (K^i - k)^+ q_t^{j,i} \ . \tag{4}$$

*with $j = \max\{j : T_j \leq \tau\}$.*

Because (2) is affine in $\Sigma$ we may write

$$A_j(\Sigma)\, q \ = \ q + G_j(q)\, \Sigma, \qquad G_j(q) \in \mathbb{R}^{(N+2)\times N}, \tag{5}$$

so that $G_j(q)$ collects the terms carrying $\Sigma$. The map $G_j$ is linear in $q$ and $G_j(q)\Sigma$ is linear in each argument separately; we use both readings. The following elementary observation is used constantly.

**Lemma 2.2** (Factorisation)**.** *Column $k$ of $G_j(q)$ involves $q$ only through $q^k$. Consequently*

$$G_j(q) \ = \ \widehat{G}_j \,\mathrm{diag}(q^1, \dots, q^N), \tag{6}$$

*where $\widehat{G}_j \in \mathbb{R}^{(N+2)\times N}$ has $k$-th column $\Delta_j\big(-\iota^{k,-}e_{k-1} + (\iota^{k,-} + \iota^{k,+})e_k - \iota^{k,+}e_{k+1}\big)$. Moreover $\widehat{G}_j$ has rank $N$, hence is injective.*

*Proof.* Immediate from (2), whose terms carrying $\Sigma^k$ are all proportional to $q^k$. For injectivity, suppose $\widehat{G}_j c = 0$ and let $k$ be the smallest index with $c_k \neq 0$. Row $k-1$ receives a contribution from column $k'$ only for $k' \in \{k-2, k-1, k\}$, and the first two carry $c_{k'} = 0$. So row $k-1$ reads $-\Delta_j \iota^{k,-} c_k = 0$, forcing $c_k = 0$, a contradiction. □

Note that $\widehat{G}_j$ depends only on the grid and on $\Delta_j$, not on the state; the state enters (6) only through the diagonal factor.

**Proposition 2.3** (The code-book)**.** *Fix the front density $q^1$, a probability vector on the grid with mean $S_0$, and let $j = 2, \dots, M$.*

*(i) Given densities $q^{j-1}$ and $q^j$ with $q^j > 0$, the discrete local variance is uniquely determined by the linear system*

$$G_j(q^j)\, \Sigma^j \ = \ q^{j-1} - q^j. \tag{7}$$

*(ii) Conversely, given $\Sigma^j > 0$, the vector $q^j := A_j(\Sigma^j)^{-1} q^{j-1}$ is again a probability vector with mean $S_0$, and the gridded prices $(\pi^i)^\top q^j$ admit no static arbitrage.*

*(iii) A gridded surface is free of static arbitrage, butterflies and calendar spreads together, if and only if it arises as in* (ii) *from some $\Sigma \geq 0$, and it lies in the* interior *of the arbitrage-free set if and only if $\Sigma > 0$.*

*Proof.* (i) is (5), with uniqueness from Lemma 2.2. For (ii), $A_j(\Sigma^j)^{-1}$ is non-negative with unit column sums and inherits the martingale property [8], attributed there to Andreasen and Huge [1]; non-negativity excludes butterfly arbitrage, and $q^j$ is then a mean-preserving spread of $q^{j-1}$, so Jensen's inequality excludes calendar arbitrage. (iii) is [8], after [3]: absence of arbitrage is equivalent to non-negative Gamma and Backward-Theta at every node, and $\Sigma$ is their ratio. Both conditions are nodewise, and $\Sigma^{j,i} = 0$ exactly when Backward-Theta vanishes. □

What everything else rests on is that the admissible set is a product of half-lines, one per node, so it has no faces in the sense that a polytope does, and $\lambda = \log \Sigma$ maps its interior diffeomorphically onto $\mathbb{R}^{N(M-1)}$.

### 2.2 Two conservation laws

The whole structure of the drift restriction follows from one observation about the left null space of $G_j$. Both identities below are used by [8], where they hold by construction for the backward operator. What is not there, because nothing in that paper needs it, is that $\operatorname{span}\{\mathbf{1}, K\}$ is the *entire* left null space. That is a rank statement, and it is what makes the drift system of Section 3 square.

**Lemma 2.4** (Left null space)**.** *For every $q$ and every $j$,*

$$\mathbf{1}^\top G_j(q) = 0 \qquad \text{and} \qquad K^\top G_j(q) = 0,$$

*where $\mathbf{1} = (1, \dots, 1)$ and $K = (K^0, \dots, K^{N+1})$. Consequently $A_j(\Sigma)$ conserves both total mass and the spot, for any $\Sigma$. Combined with Lemma 2.2, the left null space of $\widehat{G}_j$ is exactly* $\operatorname{span}\{\mathbf{1}, K\}$.

*Proof.* Fix an interior index $k$ and collect the three entries of $G_j(q)$ in column $k$: from (2) they are $\Delta_j(\iota^{k,-} + \iota^{k,+})q^k$ in row $k$, $-\Delta_j \iota^{k,+} q^k$ in row $k+1$, and $-\Delta_j \iota^{k,-} q^k$ in row $k-1$. Their sum is 0, giving the first identity. Weighting by the node values,

$$\begin{aligned}\Delta_j q^k \Big[K^k\big(\iota^{k,-} + \iota^{k,+}\big) - K^{k+1}\iota^{k,+} - K^{k-1}\iota^{k,-}\Big] &= \Delta_j q^k \Big[\iota^{k,-}(K^k - K^{k-1}) - \iota^{k,+}(K^{k+1} - K^k)\Big] \\ &= \Delta_j q^k \big[\omega^k - \omega^k\big] = 0\end{aligned}$$

by (1), giving the second. The conservation statements follow from (5): $\mathbf{1}^\top A_j(\Sigma) q = \mathbf{1}^\top q$ and $K^\top A_j(\Sigma) q = K^\top q$. Since $\widehat{G}_j$ has rank $N$ and lives in $\mathbb{R}^{(N+2)\times N}$, its left null space is two-dimensional, hence exactly $\operatorname{span}\{\mathbf{1}, K\}$. □

*Remark* 2.5 (Counting). For $q > 0$, $G_j(q)$ has $N+2$ rows and $N$ columns and rank $N$. Hence the $N+2$ martingale conditions per expiry contain exactly two redundancies, conservation of mass and of the spot, and the remaining $N$ conditions match the $N$ unknowns of $\Sigma^j$. The system is square on the quotient.

## 3 The drift restriction

**Theorem 3.1** (Martingale property)**.** *Assume that there is a measure $\mathbb{Q}$ such that each $q_t^{j,i}$ is a martingale. Then $X_t$ and all options written on it are martingales.*

Under $\mathbb{Q}$ every gridded instrument price is a martingale in $t$. Since the gridded payoffs $\pi^i$ together with the unit payoff span $\mathbb{R}^{N+2}$, and $\mathbf{1}^\top q^j = 1$, this is equivalent to each density vector being a martingale,

$$\mathrm{d}q_t^j = \psi_t^j \, \mathrm{d}W_t, \qquad \psi_t^j \in \mathbb{R}^{(N+2)\times d}. \tag{8}$$

This is the only assumption behind the drift restriction, and it carries no regularity. Each $q^j$ is a fixed linear combination of traded butterflies, so (8) is the martingale representation theorem applied to traded prices, and nothing has to commute with a stochastic integral. This is the substantive advantage of a discrete code-book over its continuum analogue, where the corresponding step requires differentiating a stochastic field in strike and maturity and must be licensed by standing regularity conditions. Theorem 3.2 below needs only $q^j > 0$.

**Theorem 3.2** (Drift restriction)**.** *Let* $\mathrm{d}\Sigma_t^j = \alpha_t^j \mathrm{d}t + \beta_t^j \mathrm{d}W_t$ *and* $q^j > 0$*. Under Assumption 3.1, for each* $j = 2, \dots, M$*,*

$$A_j(\Sigma^j)\,\psi^j \;=\; \psi^{j-1} - G_j(q^j)\,\beta^j, \tag{9}$$

*marched forward in* $j$ *from* $\psi^1$*, the volatility of the front density, which is given by its chart (Section 4). With the relative density volatility*

$$\eta^{j,i} \;:=\; \frac{\psi^{j,i}}{q^{j,i}}, \qquad i = 0, \dots, N+1, \tag{10}$$

*the drift is determined node by node,*

$$\alpha^{j,i} \;=\; -\langle \eta^{j,i}, \beta^{j,i} \rangle \qquad \Longleftrightarrow \qquad \mathrm{d}\lambda^{j,i} = \Big[ -\langle \eta^{j,i}, b^{j,i} \rangle - \tfrac{1}{2} |b^{j,i}|^2 \Big] \mathrm{d}t + b^{j,i}\,\mathrm{d}W_t, \tag{11}$$

*where* $\lambda = \log \Sigma$ *and* $b = \beta / \Sigma$*.*

*Proof.* By Lemma 2.2, $G_j(q)\Sigma = \widehat{G}_j(\Sigma \circ q)$, where $\Sigma \circ q := (\Sigma^i q^i)_{i=1}^N$ is the entrywise product on the interior nodes. The code-book therefore reads

$$q_t^j + \widehat{G}_j\, p_t^j = q_t^{j-1}, \qquad p^{j,i} := \Sigma^{j,i} q^{j,i}, \quad i = 1, \dots, N,$$

with $\widehat{G}_j$ constant. Itô's product rule for the scalar processes $\Sigma^{j,i}$ and $q^{j,i}$ gives

$$\mathrm{d}p^{j,i} = \Big[ q^{j,i} \alpha^{j,i} + \langle \psi^{j,i}, \beta^{j,i} \rangle \Big] \mathrm{d}t + \big( \Sigma^{j,i} \psi^{j,i} + q^{j,i} \beta^{j,i} \big) \mathrm{d}W_t, \tag{12}$$

with $\psi^{j,i}, \beta^{j,i} \in \mathbb{R}^{1 \times d}$ the rows at node $i$. Hence $\mathrm{d}q^j + \widehat{G}_j\,\mathrm{d}p^j = \mathrm{d}q^{j-1}$, and by uniqueness of the decomposition of a continuous semimartingale the $\mathrm{d}W$ and $\mathrm{d}t$ terms match separately. The $\mathrm{d}W$ terms give $\psi^j + \widehat{G}_j \big( \Sigma^j \circ \psi^j + q^j \circ \beta^j \big) = \psi^{j-1}$, factor by factor, which is (9) by (5). Since $q^j$ and $q^{j-1}$ are driftless by Assumption 3.1, the $\mathrm{d}t$ terms give

$$\widehat{G}_j \Big( q^j \circ \alpha^j + \textstyle\sum_f \psi^{j,\cdot,f} \circ \beta^{j,\cdot,f} \Big) = 0.$$

Since $\widehat{G}_j$ is injective (Lemma 2.2) the bracket vanishes, $q^{j,i}\alpha^{j,i} + \langle \psi^{j,i}, \beta^{j,i} \rangle = 0$ node by node, which is (11) after dividing by $q^{j,i} > 0$. The logarithmic form is Itô applied to $\lambda = \log \Sigma$, using $\alpha / \Sigma = -\langle \eta, b \rangle$. □

Only the interior components of $\eta$ enter (11), since $\Sigma$ lives on the interior nodes; the boundary components are needed in Section 5.

Equation (9) is not ours. Compared with the pricing recursion $A_j(\Sigma^j) q^j = q^{j-1}$ it is *the same operator with one added source*, the structure of Carmona and Nadtochiy [9]. The drift restriction above is theirs in the continuum, and Wissel's [19] on this lattice; we claim no priority for either.

What the fixed lattice changes is threefold. Their restriction carries a term in the derivative of the code-book along time to maturity; a lattice of fixed expiry dates has no such derivative, and here that term is absent. The relation is algebraic rather than differential, so neither the high-order regularity their statement requires nor the singular behaviour as time to maturity vanishes that it is built to accommodate can arise here, the front slab having positive width. And the discrete system is over-determined, $N + 2$ conditions on $N$ unknowns, which is what the counting of Remark 2.5 squares.

Two features of (11) are worth drawing out. The first is its shape. It is the exact analogue of the HJM restriction $\alpha = \sigma \int \sigma$: the drift of the log discrete local variance is minus the covariance of its own volatility with the volatility of the log density, plus an Itô convexity term. Where HJM's derived object is an integral of the volatility across maturities, ours is $\eta$, obtained from $b$ by the single forward recursion (9). The second is the cost. Once $\eta$ is known no matrix is inverted: no Jacobian of the pricing map is ever formed and no rank condition on it appears. And since $\psi$ is affine in $\beta$ and $\alpha$ is bilinear in $(\beta, \psi)$, the drift is a quadratic polynomial in $\beta$, as in HJM; Theorem 5.4 is what rules out the corresponding finite-time explosion.

The driftless point is the classical model: if $\beta \equiv 0$ and $\psi^1 = 0$ then $\psi \equiv 0$ and $\alpha \equiv 0$, so the discrete local variance surface is frozen, and with $\psi^1 \neq 0$ the densities still move, driven by the front alone, with $\alpha \equiv 0$ persisting.

**Proposition 3.3** (Spot consistency)**.** *The spot is $S_t = K^\top q_t^j$, the same for every $j$ by Lemma 2.4. Since $\mathbf{1}^\top A_j = \mathbf{1}^\top$, $K^\top A_j = K^\top$ and $\mathbf{1}^\top \widehat{G}_j = K^\top \widehat{G}_j = 0$, equation (9) gives*

$$\mathbf{1}^\top \psi^j = \mathbf{1}^\top \psi^1 = 0 \qquad \text{and} \qquad K^\top \psi^j = K^\top \psi^1 = \sigma_t^S S_t \quad \text{for every } j.$$

*So the spot's volatility is the single row $K^\top \psi^1 \in \mathbb{R}^{1\times d}$, it is automatically the same at every expiry, and no further condition links the spot to the surface. Taking $\psi^1 = 0$ freezes the spot and recovers $K^\top \psi^j = 0$.*

Proposition 3.3 is the discrete form of Carmona and Nadtochiy's spot consistency condition: the spot and its volatility are both *read off* the surface rather than modelled, $S = K^\top q^j$ and $\sigma^S S = K^\top \psi^j$ for any $j$. If $K^0 = 0$, the read-off is the zero-strike call, $S_t = C_t(0, T_j)$ for every $j$, as in [19] and in the multi-strike model of [15].

The drift restriction is used below through the following chain, which is what a model actually evaluates:

$$(\psi^1, b) \xrightarrow{(9)} \psi^2, \ldots, \psi^M \xrightarrow{\eta = \psi/q} \alpha, \tag{13}$$

a forward march of $d$ tridiagonal solves per expiry, $O(NMd)$ in total. The volatility structure is the only input; the drift is computed, never fitted.

# 4 The spot and the front expiry

The spot is not a state variable. By Proposition 3.3 it is the read-off $S_t = K^\top q_t^j$, the same scalar at every expiry. What has to be specified is the front density $q^1$, the initial condition of the recursion of Proposition 2.3; everything at $j \geq 2$ follows from it and from the volatility structure.

## 4.1 The front chart

The front density must be strictly positive and sum to one. Both are automatic in the chart

$$q^{1,i} = \frac{e^{z^i}}{\sum_k e^{z^k}}, \qquad \mathrm{d}z = m\,\mathrm{d}t + \sum_{f=1}^{d} \nu^f\,\mathrm{d}W^f, \tag{14}$$

with $z \in \mathbb{R}^{N+2}$ unconstrained. The chart is invariant under $z \mapsto z + c\mathbf{1}$, so a shift of $\nu$ along $\mathbf{1}$ leaves $q^1$ unchanged and only the $q^1$-centred part of $\nu$ is identified; write $\mathrm{dev}^f := \nu^f - (q^1)^\top \nu^f\,\mathbf{1}$

for it. The input $\nu$ is free, and plays the role $b$ plays at $j \geq 2$. Itô's formula for a function of a vector process shows that each $q^{1,i}$ is a martingale exactly when

$$m^i = -\tfrac{1}{2} \sum_f \big(\mathrm{dev}^{i,f}\big)^2 \tag{15}$$

up to a multiple of $\mathbf{1}$. Furthermore we obtain

$$\psi^{1,\cdot,f} = q^1 \circ \mathrm{dev}^f, \qquad \eta^{1,\cdot,f} = \mathrm{dev}^f, \qquad \mathbf{1}^\top \psi^{1,\cdot,f} = 0, \tag{16}$$

with $\circ$ the entrywise product. The same computation appears in [11] and as the stochastic replicator equation [13].

*Remark* 4.1 (Resolution at the front expiry). Among grid measures with mean $S$ the least dispersed is the pair of nodes straddling it, so

$$\mathrm{Var}_{q^1}(K) \;\geq\; (S - K^i)(K^{i+1} - S), \qquad K^i \leq S < K^{i+1}, \tag{17}$$

which is at most $\frac{1}{4}h^2$. The front implied volatility therefore has the floor

$$\frac{\sqrt{(S - K^i)(K^{i+1} - S)}}{S\sqrt{T_1 - t}} \;\leq\; \frac{h}{2S\sqrt{T_1 - t}},$$

largest at the midpoint of a cell and zero when the spot sits on a node, which is exactly what Proposition 4.2 forces at $T_1$.

### 4.2 Terminal settlement: pinning the spot to a strike

The chart constrains $q^1$ at no time $t < T_1$ beyond positivity and mass, so it imposes no terminal condition. Wissel's front parametrisation fixes the terminal state whatever the coefficients [19]; here the terminal state is a modelling choice. We take the front volatility proportional to the distance from the spot, in one factor,

$$\eta_t^{1,i} = \frac{c_t\,(K^i - S_t)}{h}, \qquad c_t = a_1 \tau^{-p} \text{ near } T_1, \quad \tau = T_1 - t,\ a_1 > 0. \tag{18}$$

Only the behaviour near $T_1$ matters: on $[0, T_1 - \varepsilon]$ the front volatility may be any bounded, locally Lipschitz choice, and (18) need only hold on $[T_1 - \varepsilon, T_1)$. It is already centred, since $\sum_i q^{1,i}(K^i - S_t) = 0$. Its coefficients are polynomial in $q^1$ and bounded on $[0, T]$ for every $T < T_1$, so there is a unique strong solution on $[0, T_1)$,

$$q_t^{1,i} = q_0^{1,i} \exp\Big( \int_0^t \eta^{1,i}\,\mathrm{d}W - \tfrac{1}{2}\int_0^t (\eta^{1,i})^2 \mathrm{d}s \Big), \qquad \|\eta_t^1\|_\infty \leq c_t\, w/h, \quad w := K^{N+1} - K^0. \tag{19}$$

**Proposition 4.2** (The front collapses onto a node, and settles exactly). *If $p \geq \frac{1}{2}$ then $q_{T_1}^1 := \lim_{t \uparrow T_1} q_t^1 = e_k$ for a random index $k$, almost surely, so $S_{T_1} = K^k$ and*

$$\sum_i q_{T_1}^{1,i} (K^i - K^n)^+ = (S_{T_1} - K^n)^+ \qquad \textit{for every } n. \tag{20}$$

*If $p < 1$ then Assumption 5.5 holds on $[0, T_1]$. The window $p \in [\frac{1}{2}, 1)$ gives both.*

*Proof.* Each $q^{1,i}$ is a bounded martingale, so $q^1_t \to q^1_{T_1}$ almost surely, and the variance $V_t := \sum_i q^{1,i}_t (K^i - S_t)^2$ converges to that of $q^1_{T_1}$. By (16) and (18) the volatility of $S = K^\top q^1$ is $\sum_i K^i q^{1,i} \eta^{1,i} = c_t V_t / h$ on $[T_1 - \varepsilon, T_1)$. Since $S$ is a martingale with values in $[K^0, K^{N+1}]$,

$$\mathbb{E} \int_{T_1-\varepsilon}^{T_1} \frac{c_t^2 V_t^2}{h^2} \, \mathrm{d}t = \mathbb{E}\big[(S_{T_1} - S_{T_1-\varepsilon})^2\big] \le w^2,$$

so the integral is finite almost surely. On the event $\{V_{T_1} > 0\}$ we would have $V_t \ge V_{T_1}/2$ near $T_1$, and the integral would diverge, because $\int^{T_1} c_t^2 \, \mathrm{d}t = a_1^2 \int^{T_1} \tau^{-2p} \, \mathrm{d}t = \infty$ for $2p \ge 1$. Hence $V_{T_1} = 0$ almost surely: $q^1_{T_1}$ is a probability vector on distinct points with zero variance, that is $e_k$ with $K^k = S_{T_1}$, and (20) reads $(K^k - K^n)^+ = (S_{T_1} - K^n)^+$. Finally $\eta^1$ is bounded on $[0, T_1 - \varepsilon]$ and $\int_{T_1-\varepsilon}^{T_1} \|\eta^1_t\|_\infty \mathrm{d}t \le (w/h) \int_{T_1-\varepsilon}^{T_1} c_t \, \mathrm{d}t < \infty$ when $p < 1$, by (19). □

The terminal spot is therefore a grid point, as in [19]; the difference is that here it is chosen, by (18), and there it is forced by the parametrisation.

*Remark* 4.3 (What the lattice could carry, and does not). At expiry each butterfly price $q^{j,i}$ is worth its payoff: for $S_{T_j} \in [K^k, K^{k+1}]$ the second difference defining $q^{j,i}$ telescopes to $\varpi(S_{T_j})$, the barycentric weights of the spot on the two nodes straddling it. Since $u \mapsto (u - K^n)^+$ is affine between nodes,

$$\sum_i \varpi^i(s) \, (K^i - K^n)^+ \;=\; (s - K^n)^+ \qquad \text{for every } s \in [K^0, K^{N+1}] \text{ and every model strike } K^n, \tag{21}$$

so $\varpi(s)$ is an admissible terminal state that settles every gridded contract exactly, also for $s$ strictly inside a cell. The chart does not force the pinning; the choice (18) does, because its terminal variance must vanish. A front volatility that drives $q^1$ to $\varpi(s)$ with $s$ inside a cell is left as future work (Section 7).

## 5 The wing constant, and existence

The log discrete local variance $\lambda = (\lambda^2, \dots, \lambda^M)$, $\lambda^j = \log \Sigma^j$, ranges over all of $\mathbb{R}^{N(M-1)}$, so positivity is automatic; the front density has its own chart (14). By Theorem 3.2 the model is the stochastic differential equation

$$\mathrm{d}\lambda_t = \Big[ -\langle \eta_t, b_t \rangle - \tfrac{1}{2} |b_t|^2 \Big] \mathrm{d}t + b_t \, \mathrm{d}W_t, \tag{22}$$

whose only non-elementary coefficient is $\eta = \psi / q$ from (9) and (10). Its denominator is small in the wings, and the question is whether $\eta$ can blow up there.

The argument has the structure of Wissel's proof of his Theorem 3.5 [19]: log coordinates, induction over expiries, and local existence followed by a bound that rules out explosion. Its one non-standard ingredient is an a priori bound, independent of the state, on the relative volatility of the calendar spread at the second expiry. Wissel could not obtain this bound in general and imposes his condition (3.24) instead (his Remark 1 following the proof, p. 18); in our coordinates the quantity is $\eta + b$. We obtain it from an exact identity for the volatility recursion and a known inequality for inverses of M-matrices.

### 5.1 The volatility recursion

Write $P_j := A_j(\Sigma^j)^{-1}$, so that $q^j = P_j q^{j-1}$, and $P$ when $j$ is fixed.

**Corollary 5.1** (Positivity, unit column sums, absorbing ends)**.** *$P \geq 0$ entrywise, $\mathbf{1}^\top P = \mathbf{1}^\top$, so $P_{im} \in [0,1]$, and the outer columns of $P$ are $e_0$ and $e_{N+1}$.*

*Proof.* By (2), $A_j$ has non-positive off-diagonal entries, and in each column the diagonal entry exceeds the sum of the moduli of the off-diagonal ones by exactly one; so $A_j$ is a non-singular M-matrix and $P \geq 0$ [2]. The column sums follow from $\mathbf{1}^\top A_j = \mathbf{1}^\top$ (Lemma 2.4), and the outer columns of $A_j$ are $e_0$ and $e_{N+1}$. The first two statements are also Theorem 3.3 of [8], attributed there to [1]. □

That is all that is needed about $P$ itself. What we want next is a recursion for $\eta$ alone: one step per expiry, each adding a controlled amount, so that the bound telescopes along the ladder. Dividing (9) by $q^j$ gives

$$\eta^{j,i} \;=\; \frac{(P\psi^{j-1})_i}{q^{j,i}} \;-\; \frac{\big(P\,G_j(q^j)\beta^j\big)_i}{q^{j,i}}, \tag{23}$$

the first term carrying $\eta^{j-1}$ and the second carrying $b^j$.

The first is an average. Writing $\psi^{j-1} = q^{j-1} \circ \eta^{j-1}$, its weights $P_{ik}q^{j-1,k}/q^{j,i}$ are non-negative because $P \geq 0$ and sum to one because $q^j = Pq^{j-1}$, so it contributes at most $\|\eta^{j-1}\|_\infty$; they are the matrix $\mathcal{A}$ of (28).

The source is the difficulty. Since $G_j$ is a second difference, row $i$ carries the neighbouring densities $q^{j,i\pm1}$ as well as $q^{j,i}$, so dividing by $q^{j,i}$ leaves the ratios

$$\frac{q^{j,i-1}}{q^{j,i}} \qquad \text{and} \qquad \frac{q^{j,i+1}}{q^{j,i}}, \tag{24}$$

which are unbounded in the wings, where the density decays fastest; nothing in (9) bounds them.

They are, however, an artefact of writing the row out. Since $\mathbf{1}^\top G_j(q) = 0$ by Lemma 2.4, row $i$ of $G_j(q)\Sigma^j$ is the net outflow of node $i$, and by (5) that net is

$$G_j(q^j)\,\Sigma^j \;=\; A_j(\Sigma^j)q^j - q^j \;=\; q^{j-1} - q^j, \tag{25}$$

which carries no ratios at all: the individual fluxes are wing quantities, their net is not. For $b^j \equiv c$ this settles the matter, since $\beta^j = c\,\Sigma^j$ and $Pq^{j-1} = q^j$ give

$$\frac{\big(P\,G_j(q^j)\beta^j\big)_i}{q^{j,i}} = c\,\frac{\big(q^j - Pq^j\big)_i}{q^{j,i}} = c\Big(1 - \frac{(Pq^j)_i}{q^{j,i}}\Big),$$

the two ratios of (24) having collapsed into the single propagation ratio $(Pq^j)_i/q^{j,i}$. A general $b$ only modulates the density. Let $E \in \mathbb{R}^{N\times(N+2)}$ restrict to the interior nodes, so $EE^\top = I_N$, and note that $A_j = I + \widehat{G}_j\,\mathrm{diag}(\Sigma^j)E$. Writing $E^\top(q \circ b)$ for the vector equal to $q^i b^i$ at the interior nodes and zero at the outer ones, (6) and the commuting of diagonal matrices give

$$G_j(q)\beta^j = \widehat{G}_j\,\mathrm{diag}(\Sigma^j)E\,E^\top(q \circ b^j) = \big(A_j(\Sigma^j) - I\big)E^\top(q \circ b^j), \tag{26}$$

again the operator's own increment, applied to $E^\top(q \circ b^j)$ in place of $q$. The proposition is what this amounts to.

**Proposition 5.2.** *With $b^j = \beta^j/\Sigma^j$ and $\circ$ the entrywise product,*

$$\psi^j \;=\; P\,\psi^{j-1} \;-\; (I-P)\,E^\top\big(q^j \circ b^j\big), \tag{27}$$

*equivalently, with the entrywise non-negative*

$$\mathcal{A} := \operatorname{diag}(q^j)^{-1} P \operatorname{diag}(q^{j-1}), \qquad \mathcal{B} := \operatorname{diag}(q^j)^{-1} P\,E^\top \operatorname{diag}(Eq^j),$$

$$\eta^j \;=\; \mathcal{A}\,\eta^{j-1} \;+\; \big(\mathcal{B} - E^\top\big) b^j, \tag{28}$$

*where $\mathcal{A}\mathbf{1} = \mathbf{1}$, so the transport is an averaging operator, and the $i$-th row of $\mathcal{B}$ sums to $(Pq^j)_i/q^{j,i} - \mathbf{1}_{\{i\in\{0,N+1\}\}}$.*

*Proof.* By (9), $\psi^j = P\psi^{j-1} - P\,G_j(q^j)\beta^j$. Applying $P$ to (26) at $q = q^j$ and using $PA_j = I$ gives $P\,G_j(q^j)\beta^j = (I-P)E^\top(q^j \circ b^j)$, which is (27). Dividing it by $q^{j,i}$ gives (28): the rows of $\mathcal{A}$ sum to $(Pq^{j-1})_i/q^{j,i} = 1$, and for $\mathcal{B}$ the terms $k \in \{0, N+1\}$ missing from $(Pq^j)_i$ vanish for interior $i$ and equal $q^{j,i}$ at $k = i$ for outer $i$, by Corollary 5.1. □

So $\eta^j$ is an average of $\eta^{j-1}$ plus a source in $b^j$, and the march across expiries is additive.

## 5.2 The wing constant

That single surviving ratio, $(Pq^j)_i/q^{j,i}$, is what has to be bounded. Since $q^j = Pq^{j-1}$ it is

$$\frac{(Pq^j)_i}{q^{j,i}} = \frac{(P^2q^{j-1})_i}{(Pq^{j-1})_i}, \tag{29}$$

so what is needed is how much larger than $P$ the square $P^2$ can be entrywise; and since both sides are linear and non-negative in $q^{j-1}$, the case $q^{j-1} = e_m$ shows that an entrywise bound is not merely sufficient but necessary. Write

$$R^{j,i} := \mathbf{1}_{\{1\le i\le N\}} + \frac{(Pq^j)_i}{q^{j,i}} - \mathbf{1}_{\{i\in\{0,N+1\}\}}, \qquad \Theta_j := \max_i R^{j,i}, \qquad \Xi_j := \sum_{r=1}^{N} P_{rr} \le N, \tag{30}$$

and $H_j := \max_i|\eta^{j,i}|$. Here $R^{j,i}$ is the row bound of (28) at node $i$, and $\Theta_j$, the *wing constant*, is its size across the slab. The theorem shows $\Theta_j$ is bounded by the grid alone, uniformly in the state; the bound rests on a single inequality for inverses of M-matrices.

**Lemma 5.3** (Squared propagator)**.** *For all nodes $i, r, m$,*

$$P_{ir}\,P_{rm} \;\le\; P_{im}\,P_{rr}, \tag{31}$$

*and consequently, for all $i, m$,*

$$\big(P^2\big)_{im} \le \begin{cases} \Xi_j\,P_{im}, & 1 \le i \le N, \\ (1+\Xi_j)\,P_{im}, & i \in \{0, N+1\}. \end{cases} \tag{32}$$

*Proof.* $P$ is the inverse of an M-matrix by Corollary 5.1, and every inverse M-matrix is a path product matrix, $a_{ij}a_{jk} \le a_{ik}a_{jj}$ for all $i, j, k$ [12, Theorem 4.5]; (31) is the case $(i, j, k) = (i, r, m)$. Summing it over $r$ gives (32): an outer $r$ contributes only if $r = i$, since column $r$ of $P$ is $e_r$, and then contributes $P_{im}$. □

**Theorem 5.4.** *For every $\Sigma \in (0,\infty)^N$,*

$$H_j \;\le\; H_{j-1} + \Theta_j \|b^j\|_\infty, \tag{33}$$

*and*

$$\Theta_j \le 1 + \Xi_j \le N+1, \qquad \Theta^* := \sup_{\Sigma} \sum_{j=2}^{M} \Theta_j \le (M-1)(N+1). \tag{34}$$

*Proof.* Since $q^j = Pq^{j-1} \ge 0$, Lemma 5.3 gives $(Pq^j)_i = \sum_m (P^2)_{im} q^{j-1,m} \le \Xi_j\, q^{j,i}$ for interior $i$ and $\le (1+\Xi_j) q^{j,i}$ for outer $i$, so $R^{j,i} \le 1 + \Xi_j$ in both cases. Finally, $\mathcal{A} \ge 0$ with $\mathcal{A}\mathbf{1} = \mathbf{1}$ gives $\|\mathcal{A}\eta^{j-1}\|_\infty \le H_{j-1}$, while row $i$ of $(\mathcal{B} - E^\top) b^j$ has modulus at most $\|b^j\|_\infty R^{j,i}$ by the row sums of Proposition 5.2. □

## 5.3 Existence

Iterating (33) from the front slice,

$$\sup_{j,i} |\eta_t^{j,i}| \;\le\; \|\eta_t^1\|_\infty + (M-1)(N+1)\, \|b\|_\infty, \tag{35}$$

and iterating (27) shows that $\eta^j$ is affine in $\eta^1$, with coefficients depending on $\lambda$ only through $P_2, \dots, P_j$ and $q^2, \dots, q^j$. The only $t$ on the right of (35) is inside $\eta_t^1$, which enters linearly, so it has to be integrable in time rather than bounded.

**Assumption 5.5** (Integrable front volatility)**.** For a horizon $T \le T_1$,

$$\int_0^T \|\eta_t^1\|_\infty \, \mathrm{d}t \;<\; \infty \qquad \text{almost surely.} \tag{36}$$

**Theorem 5.6** (Existence)**.** *Let the front volatility $\nu = \nu(t,z)$ in (14) and the volatility $b = b(t,z,\lambda)$ be locally Lipschitz in $(t,z,\lambda)$ on $[0,T_1)$, with $|b| \le B$, and $\nu$ bounded on $[0,T]$ for every $T < T_1$.*

*(i) For every $T < T_1$ the front chart (14)–(15) together with (22) has a unique strong solution $(z,\lambda)$ on $[0,T]$, and*

$$0 \;<\; \inf_{t \le T} \min_{j,i} \Sigma_t^{j,i} \;\le\; \sup_{t \le T} \max_{j,i} \Sigma_t^{j,i} \;<\; \infty \qquad \textit{almost surely.} \tag{37}$$

*(ii) If moreover Assumption 5.5 holds with $T = T_1$, then $\lambda_t$ has a finite limit $\lambda_{T_1}$ as $t \uparrow T_1$, and (37) holds with $T = T_1$.*

*No condition is imposed on $\Sigma$, on the wings, or on the number of factors.*

*Proof.* (i) Fix $T < T_1$ and let $B_\nu$ bound $|\nu|$ on $[0,T]$. The coefficients are locally Lipschitz: $q^1 = \mathrm{softmax}(z)$ and $\Sigma = e^\lambda$ are positive, so each $A_j$ is a non-singular M-matrix and everything except $\nu$ and $b$ is smooth in $(z,\lambda)$. Adding time as a state variable to cover the time dependence, [14, Theorem V.38, p. 303] gives a unique solution up to an explosion time $\zeta$, with $\limsup_{t\to\zeta} \|(z_t,\lambda_t)\| = \infty$ on $\{\zeta < \infty\}$. The drift of $\lambda$ is bounded by (35), since $\|\eta^1\|_\infty = \|\mathrm{dev}\|_\infty \le 2B_\nu$, and the drift $m$ of $z$ is bounded by $2B_\nu^2$; let $D$ bound both. On $[0, \zeta \wedge T)$ the drift part of $(z,\lambda)$ is therefore at most $DT$, and the noise part, having integrand bounded by $B_\nu + B$, is a martingale bounded in $L^2$, hence convergent as $t \to \zeta \wedge T$. So $(z,\lambda)$ stays bounded

on $[0, \zeta \wedge T)$, which rules out $\zeta \le T$. The solution exists on $[0,T]$, $\lambda$ is bounded there, and (37) follows from $\Sigma = e^{\lambda}$.

(ii) By (i) the solution exists on $[0,T_1)$. By (35) the drift of $\lambda$ is bounded by $g(t) := \big(\|\eta_t^1\|_\infty + (M-1)(N+1)B\big)B + \frac{1}{2}B^2$, which is integrable on $[0,T_1)$ by Assumption 5.5, so the drift part of $\lambda$ converges as $t \uparrow T_1$. The noise part $\int_0^t b\,\mathrm{d}W$ has integrand bounded by $B$, so it is a martingale bounded in $L^2$ on $[0,T_1)$ and converges as well. Hence $\lambda_t$ has a finite limit, $\lambda$ is bounded on $[0,T_1]$, and (37) holds with $T = T_1$. Nothing is claimed for $z$, which may run off at $T_1$ as the front collapses. □

The settling front (18) is bounded on every $[0,T]$ with $T < T_1$ by (19), and satisfies Assumption 5.5 with $T = T_1$ for $p < 1$ (Proposition 4.2). So for $p \in [\frac{1}{2}, 1)$ the code-book reaches its first expiry with $\Sigma$ bounded away from 0 and $\infty$ and every gridded contract settling at its payoff.

### 5.4 Beyond the first expiry

At $T_1$ the front expires and the next expiry becomes the front. The settling of Proposition 4.2 makes this switch clean: the spot sits on a node, and the new front density is a column of the propagator. Indeed, if $q^l_{T_l} = e_k$ then the code-book $A_{l+1}(\Sigma^{l+1})\,q^{l+1} = q^l$ gives

$$q^{l+1}_{T_l} = P_{l+1}\,e_k, \qquad P_{l+1} := A_{l+1}\big(\Sigma^{l+1}_{T_l}\big)^{-1}, \tag{38}$$

the spread of a unit mass at the settled spot by one step of the code-book. To restart the front chart we need this density to be strictly positive.

**Lemma 5.7** (The new front)**.** *Let $P = A_j(\Sigma^j)^{-1}$ with $\Sigma^j \in (0,\infty)^N$. If $k$ is an interior node, then $(Pe_k)_i > 0$ for every node $i$ and $K^\top P e_k = K^k$. If $k \in \{0, N+1\}$, then $Pe_k = e_k$.*

*Proof.* Order the nodes with the interior first. The outer columns of $A_j$ are $e_0$ and $e_{N+1}$ (Corollary 5.1), so

$$A_j = \begin{pmatrix} A_{II} & 0 \\ A_{BI} & I \end{pmatrix}, \qquad P = \begin{pmatrix} A_{II}^{-1} & 0 \\ -A_{BI}A_{II}^{-1} & I \end{pmatrix}.$$

By (2), $A_{II}$ is tridiagonal with strictly negative off-diagonal entries, hence an irreducible non-singular M-matrix, so $A_{II}^{-1} > 0$ entrywise [2]. The block $-A_{BI} \ge 0$ has a positive entry in the row of $K^0$ (column $K^1$) and in the row of $K^{N+1}$ (column $K^N$), so the outer entries of an interior column of $P$ are positive as well. The identity $K^\top P = K^\top$ follows from $K^\top A_j = K^\top$ (Lemma 2.4), and the outer columns are Corollary 5.1. □

**Corollary 5.8** (Existence on the whole ladder)**.** *Let $b$ be bounded and locally Lipschitz. On each interval $[T_{l-1}, T_l]$, $l = 1, \dots, M$, let the current front density $q^l$ follow the chart* (14) *with the settling volatility* (18)*, where $\tau = T_l - t$ and $p \in [\frac{1}{2}, 1)$. Then the model has a solution on $[0, T_M]$, unique given these choices, such that*

*(i) every gridded price is a continuous martingale up to its expiry;*

*(ii) at every $T_l$ the spot sits on a node $K^{k_l}$ and every contract expiring at $T_l$ settles at its payoff;*

*(iii) if $k_l$ is interior, the model continues on $[T_l, T_{l+1}]$ from the finite state $z = \log q^{l+1}_{T_l}$ and $\lambda^{l+2}, \dots, \lambda^M$ at $T_l$; if $k_l$ is an outer node, the spot is absorbed there: $q^j_t = e_{k_l}$ for all $j > l$ and $t \ge T_l$.*

*Proof.* By induction on $l$. On $[0, T_1]$ this is Proposition 4.2 with Theorem 5.6, whose part (ii) also gives the limit $\lambda_{T_1}$, finite. Suppose the solution exists on $[0, T_l]$ with $\lambda_{T_l}$ finite and $q^l_{T_l} = e_k$.

If $k$ is interior, (38) and Lemma 5.7 give a strictly positive new front, so $z = \log q^{l+1}_{T_l}$ is finite, and $K^\top q^{l+1}_{T_l} = K^k = S_{T_l}$, so the spot is continuous. On $[T_l, T_{l+1}]$ the ladder is that of Sections 3–5 with front $q^{l+1}$ and expiries $T_{l+1} < \cdots < T_M$, started from the $\mathcal{F}_{T_l}$-measurable state $(z, \lambda^{l+2}, \ldots, \lambda^M)$. Theorem 5.6 applies with this $\mathcal{F}_{T_l}$-measurable initial condition, the solution of [14, Theorem V.38] depending continuously on its starting point, and with Proposition 4.2 and part (ii) of that theorem gives a unique strong solution on $[T_l, T_{l+1}]$ that settles at $T_{l+1}$ with $\lambda_{T_{l+1}}$ finite. The densities $q^j$, $j > l+1$, are given by the same recursion on both sides of $T_l$ with continuous inputs, so they are continuous at $T_l$.

If $k$ is an outer node, Lemma 5.7 gives $q^j_{T_l} = e_k$ for every $j > l$. The constant process $q^j_t \equiv e_k$ satisfies every recursion, since $A_j e_k = e_k$, and is a martingale, so all prices are constant and settle at their payoffs; we freeze $\lambda$ at $\lambda_{T_l}$, on which no price then depends.

Uniqueness on each interval gives uniqueness on $[0, T_M]$. □

# 6 Finite-step simulation and numerical evidence

The drift restriction of Section 3 admits two numerically distinct implementations. The first integrates the log-DLV SDE directly. It is the closest implementation of the continuous-time model, but a finite time step is only approximately martingale preserving. The second advances the probability density itself through a finite set of symmetric branches. It is an exact one-step martingale construction on the modeled finite support, although it is not a claim about a continuous-strike or unbounded-support limit. We report both because their numerical guarantees differ.

All experiments in this section use four weekly expiries and fixed cash strikes within each weekly transition. "DLVsq" denotes the positive discrete local variance $\Sigma$. The final density-branch audit uses the discrete factor $X$ alone. Earlier development and long-horizon audits also multiply $X$ by an independent mean-one Black factor $Y$ with annual volatility 10%; the conditioning identities below apply to both, using the Black semigroup for $Y$.

## 6.1 Direct integration of the DLV diffusion

For the direct SDE experiment we work in log excess-DLV coordinates

$$z^{j,i} = \log(\Sigma^{j,i} - \Sigma_{\min}).$$

The volatility of $z$ is the rank-25 residual covariance estimated from 703 weekly historical transitions through 31 December 2024. At every substep the density volatility is obtained from (9), and the drift of $\Sigma$ is then computed from (11), including the Itô term in log coordinates. Each week is integrated by ten stochastic-Heun substeps, using the same Gaussian increment in predictor and corrector. At the weekly endpoint spot is drawn from the front density and the expired slab is removed.

The audit starts from historical state 44 (8 March 2022), uses 200 interior strikes, 100,000 paths arranged as 50,000 antithetic pairs, and follows a shrinking tower for four weeks without re-gridding. Fifty antithetic pairs are aggregated into each of 1,000 independent randomized-stratification blocks. There are 800, 600, 400, and 200 live fixed-contract option returns at the four trade dates, hence 2,000 coordinates in the joint test.

For this experiment the blockwise coordinate statistics are supplemented by a 10,000-draw shared-sign multiplier bootstrap. If $t_k$ is the usual block t-statistic, the two simultaneous statistics are

$$T_{\max} = \max_k |t_k|, \qquad T_{\mathrm{RMS}} = \left( d^{-1} \sum_{k=1}^{d} t_k^2 \right)^{1/2}.$$

The reported decision value is $p_{\mathrm{comb}} = \min\{1, 2\min(p_{\max}, p_{\mathrm{RMS}})\}$. The shared multiplier for all coordinates of a block preserves their observed cross-option dependence.

| Audit | Active | $\max \lvert t \rvert$ | RMS $t$ | $p_{\mathrm{comb}}$ | 5% decision |
|---|---|---|---|---|---|
| One-step option returns | 2,000 | 3.334 | 1.109 | 0.2746 | do not reject |
| Terminal option returns | 2,000 | 3.334 | 1.044 | 0.3210 | do not reject |
| Terminal option price changes | 2,000 | 2.934 | 1.156 | 0.3500 | do not reject |
| Integrated-DLV control | 1,200 | 1.950 | 0.878 | 1.0000 | do not reject |
| Weekly spot returns | 4 | 2.431 | 1.561 | 0.0876 | do not reject |

Table 1: Martingale diagnostics for the ten-substep Heun implementation. The 800 omitted integrated-DLV coordinates have effectively zero block variance. The primary one-step audit has separate multiplier-bootstrap values $p_{\max} = 0.1373$ and $p_{\mathrm{RMS}} = 0.2465$; its simultaneous 5% max-$t$ cutoff is 4.235.

The largest density-mass residual was $9.33 \times 10^{-14}$ and the largest conditional-forward residual was $8.93 \times 10^{-14}$. The minimum density was zero, due to finite-support point masses rather than a negative density. The simulation took 72.1 seconds, or approximately 1,388 paths per second. These results support the discretization, but they do not turn an infinitesimal drift restriction into an exact one-week transition.

### 6.2 An exact finite-step density construction

We next construct the one-week law in density coordinates. Let $\varepsilon_b \in \{-1, +1\}^d$ be the $b$-th Rademacher branch and let $\pi_b = 2^{-d}$, with $d = 4$ in the experiment. Write $\Gamma^j$ for the state-dependent loading of density $q^j$ implied by the chosen DLVsq loading. The branch marginal is

$$q^{j,b} = q^j + \sqrt{h}\,\Gamma^j \varepsilon_b, \qquad b = 1, \ldots, 2^d. \tag{39}$$

Since $\sum_b \pi_b \varepsilon_b = 0$ and $\sum_b \pi_b \varepsilon_b \varepsilon_b^\top = I_d$,

$$\sum_b \pi_b q^{j,b} = q^j \tag{40}$$

holds exactly, before Monte Carlo sampling. DLVsq is then recovered algebraically from each $q^{j,b}$ by Proposition 2.3; it is not advanced by an Euler step or fitted by a nonlinear optimizer. A branch is accepted only if its masses are non-negative and the recovered DLVsq lies in its configured open bounds. The implementation checks mass, mean, decoder consistency, and DLV bounds at every step.

In the implemented weekly law the front density is branch invariant, $q^{1,b} = q^1$. Independently of $b$, a front node $I = i$ is drawn with probability $q_i^1$. For every surviving expiry $j \geq 2$, let $R_{ir}^{b,j}$ be the conditional probability of node $r$ given branch $b$ and front node $i$. The SANOS transition

supplies rows satisfying

$$\sum_r R_{ir}^{b,j} = 1, \qquad \sum_r x_r R_{ir}^{b,j} = x_i, \qquad \sum_i q_i^1 R_{ir}^{b,j} = q_r^{j,b}. \tag{41}$$

The first two identities make every row a martingale transition; the third recovers the branch marginal. Together with (40), they imply the tower property for every payoff represented on the finite support.

This construction is exact on the finite model grid. It is not asserted to be the unique risk-neutral law, and the particular independent sampling of $b$ and $I$ does not generate contemporaneous correlation between the raw surface branch and front-node innovation. Realized density-weighted factor covariance nevertheless matched the intended empirical covariance closely: in the independent all-state confirmation the largest relative error was 1.87% and the RMS error was 0.55%.

### 6.3 Conditioning and re-gridding

Suppose current cash spot is $s$ and the revealed front node is $x_i$, so the next $X$ spot is $sx_i$. Define the next relative support by

$$x_r'^{(i)} = \frac{x_r}{x_i}. \tag{42}$$

Then $(sx_i)x_r'^{(i)} = sx_r$ path by path. Consequently the price of a surviving fixed-cash call is identical on the inherited and re-gridded coordinates,

$$C_{t+h}^{b,i,j}(K) = \sum_r R_{ir}^{b,j}(sx_r - K)^+.$$

Averaging first over $i$ and then over $b$ gives

$$\begin{aligned}\mathbb{E}_t[C_{t+h}^j(K)] &= \sum_b \pi_b \sum_i q_i^1 \sum_r R_{ir}^{b,j}(sx_r - K)^+ \\ &= \sum_r q_r^j (sx_r - K)^+ = C_t^j(K).\end{aligned}$$

Thus re-gridding is an exact coordinate change for all surviving fixed-cash claims, not an interpolation. With an independent Black factor $Y$, the same proof conditions on $Y$ and then applies its mean-one Black semigroup. A newly issued far expiry is different: there is no pre-existing fifth-week claim whose price it must preserve.

At each weekly roll, the old 2W–4W densities become the new floating 1W–3W densities through (41) and (42). Only the new 4W DLVsq curve is supplied statistically. The retained boundary model is a 25-factor PCA followed by an unrestricted weekly VAR(1) in factor scores. It uses 703 training pairs through 2024, explains 99.9912% of the 4W curve variation, and has spectral radius 0.9089. On 155 untouched 2025 pairs its raw-coordinate RMSE is 0.5902, against 0.6526 for persistence, a 9.56% improvement. SANOS constructs the new 4W density as a static martingale extension of the inherited 3W density.

### 6.4 Martingale audit for the exact density step

Adjacent simulated paths form an antithetic pair. For pair $a = 1, \ldots, n$ and audited return coordinate $k$, let $R_{ak}$ be the pair mean, $\bar{R}$ the vector of sample means, and $\widehat{\Sigma}$ their sample

covariance. We report

$$z_k = \frac{\sqrt{n}\,\bar{R}_k}{\widehat{\sigma}_k}, \qquad J = n\bar{R}^\top \widehat{\Sigma}^+ \bar{R}. \tag{43}$$

The pseudo-inverse is formed in correlation space after discarding eigenmodes at or below 0.01. If $\nu$ modes remain, the reference value is $J \sim \chi^2_\nu$ under the joint zero-mean null. A small upper-tail $p$-value rejects zero drift; a large value means only that the observed drift is compatible with Monte Carlo noise. We use 0.05 for one look, 0.0125 for four weekly looks, and 0.003125 for 16 looks. Because the options are highly collinear and the retained modes are estimated from the sample, the exact density-tower identities and matched cash/re-grid prices are stronger structural diagnostics than an isolated asymptotic $p$-value.

The canonical X-only audit used all 891 certified historical states, 100,000 paths (50,000 independent antithetic pairs), 11 ITM call/put strikes per expiry, four one-week steps, 64-bit arithmetic, and 10% annualized log-DLVsq volatility. Each state received 56 or 57 independent pairs. The audited book contains spot and the surviving original fixed-cash claims.

| Week | Worst $\lvert z\rvert$ | Max option mean (bp) | $J$ | $\nu$ | $p$ |
|---|---|---|---|---|---|
| 1 | 1.796 | 0.497 | 10.980 | 7 | 0.139 |
| 2 | 2.143 | 0.863 | 7.228 | 7 | 0.406 |
| 3 | 1.371 | 0.696 | 5.858 | 6 | 0.439 |
| 4 | 2.057 | 1.245 | 5.506 | 5 | 0.357 |

Table 2: All-state X-only martingale audit for the surviving original fixed-cash book plus spot. Option means are basis points of initial spot. No horizon rejects the joint zero-mean null at 5%.

The deterministic checks close much more tightly than Monte Carlo sampling: the maximum density-tower residual is $6.49 \times 10^{-15}$, option-tower residual $1.67 \times 10^{-16}$, spot-mean residual $2.60 \times 10^{-14}$, four-roll decoder residual $1.77 \times 10^{-13}$, and first re-grid cash-price residual $3.33 \times 10^{-16}$.

A matched run decomposed every weekly change into the fixed-expiry surface branch and the spot-conditioning/re-grid component, using identical random draws. Direct inherited-cash and re-gridded valuations agree pathwise to at most $2.22 \times 10^{-16}$ and have identical $z/J/p$ statistics.

| Week | Fixed-expiry $p$ | Full first-member $p$ | Full pair $p$ | Max price difference |
|---|---|---|---|---|
| 1 | 0.366 | 0.684 | 0.139 | 0 |
| 2 | 0.538 | 0.163 | 0.128 | $2.22 \times 10^{-16}$ |
| 3 | 0.911 | 0.414 | 0.0208 | $2.22 \times 10^{-16}$ |
| 4 | 0.945 | 0.0808 | 0.546 | $2.22 \times 10^{-16}$ |

Table 3: Matched fixed-support and floating-coordinate decomposition. The first two $p$-value columns use one independent member of each pair; the third uses antithetic pair means.

The week-3 pair-mean value is a nominal rejection at 5%, but not at the pre-specified four-look cutoff 0.0125. It is also visibly sensitive to the correlation eigenvalue cutoff: its $p$-value is 0.0429, 0.0208, 0.0879, and 0.295 at cutoffs 0.005, 0.01, 0.02, and 0.05. Estimating covariance and the mean on opposite half-samples gives 0.214 and 0.338. No coordinate has $|z| > 2$ and the largest mean is 0.363bp. Since the direct and re-gridded path matrices are identical, this fluctuation cannot be caused by re-gridding. The fixed-expiry branch alone has maximum option means of only 0.0046, 0.0056, 0.0013, and 0.0050bp over weeks 1–4.

Earlier development audits reached the same conclusion under a nonzero Black backbone. On 21 states with 10,000 paths per state, the four weekly maximum coordinate $|z|$ values were 2.400, 3.203, 2.871, and 2.968, while the covariance-aware Fisher-combined state-test $p$-values were 0.871, 0.672, 0.852, and 0.522. A heterogeneous 10,000-path confirmation covering all 891 states gave unconditional maximum $|z|$ values 1.575, 1.970, 1.878, and 1.040 and $p$-values 0.702, 0.414, 0.067, and 0.317. Its high-dimensional state-feature-weighted values were 0.027, 0.759, 0.251, and 0.0089; they are reported because they expose the poor calibration of hard spectral cutoffs in small per-state samples, rather than being hidden behind the aggregate result.

### 6.5 Repeated rolls, speed, and finite-support behavior

A 16-week audit with the independent 10% Black backbone used 10,000 paths and all 891 states. Across all rolls the largest decoder residual was $2.35 \times 10^{-13}$, conditional mass and mean errors stayed below $6.6 \times 10^{-14}$, and DLVsq remained strictly inside $(10^{-4}, 9)$. The maximum spot $|z|$ over 16 horizons was 2.126, the maximum rolling-option $|z|$ was 2.417, and every rolling-option joint $p$-value was at least 0.0839. With the last-expiry AR boundary, median cross-expiry dispersion of $\sqrt{\text{DLVsq}}$ was 0.0459 after week 1 and 0.0569 after week 16, so the term structure did not collapse. The corresponding implementation ran 160,000 path-week steps in 14.6 seconds.

The exact branch method is materially faster than a nonlinear finite-step solve. In the heterogeneous 10,000-path benchmark it processed 40,000 path-week steps in 8.84 seconds (4,527 per second), including 7.36 seconds to construct 5,000 distinct initial kernels. This was about 23 times faster than the earlier finite-$h$ Newton transition (195 per second), but about 5.7 times slower than the local SDE transition (25,649 per second). The final X-only audit, with more extensive all-state diagnostics, processed 400,000 path-week steps in 169.8 seconds. These timings are indicative rather than a controlled hardware benchmark.

The remaining limitation is economic rather than a floating-point failure. If the lowest support node $x_0$ is revealed, then

$$\sum_r R_{0r}^{b,j} x_r = x_0, \qquad x_r \geq x_0,$$

forces $R_{00}^{b,j} = 1$. The upper endpoint is analogous. The $X$ law is then absorbing, all surviving options are intrinsic, and interior DLVsq coordinates are price-unidentified. The positive-density formula $\eta = \psi/q$ no longer applies at zero-mass nodes; any continued DLV coordinates there are an economically invisible numerical convention.

In the four-week X-only run, boundary incidence was 0.003%, 0.003%, 0.003%, and 0.006%. In the 16-week run, the fully collapsed fraction grew from 0.01% to 0.48%; the median week-16 front density retained 23.44 effective nodes, while its first percentile retained 1.15. Avoiding this absorbing tail requires dynamic support expansion, an explicit tail continuation, or a boundary state carrying a conditional tail. Re-gridding alone cannot create support beyond an extreme node while preserving the inherited option surface.

The numerical evidence therefore supports three distinct statements: the Heun implementation is a credible approximation to the continuous density drift; the symmetric density construction is exactly martingale consistent on the finite support; and weekly re-gridding preserves every surviving fixed-cash price path by path. It does not establish uniqueness, an unbounded-support model, or the absence of boundary bias at horizons much longer than those tested.

# 7 What is open

1. *An off-lattice terminal spot.* By Remark 4.3 the chart represents every $s \in [K^0, K^{N+1}]$ at expiry, and the state that does so is $\varpi(s)$, carried by the two nodes straddling $s$. Whether a front volatility exists driving $q^1$ to $\varpi(s)$ for $s$ strictly inside a cell, rather than to a single node as (18) does, is open. The obstruction is that exact settlement forces $q^1_{T_1} = \varpi(S_{T_1})$ and hence $\int_0^{T_1} (\eta^{1,i})^2 \mathrm{d}t < \infty$ for both surviving nodes, which by the argument of Proposition 4.2 they cannot both satisfy when the volatility is proportional to distance from the spot. Until it is answered the model of Proposition 4.2 pins, as [19] does, and his question of a front that does not pin the spot remains open.

2. *The continuum limit.* The constant $(M-1)(N+1)$ grows with the grid. Whether $\sup_N \Theta_j < \infty$ on a class of realistic surfaces, which is what a continuum statement would need, is open.

# References


[1] Jesper Andreasen and Brian Norsk Huge. Volatility interpolation. Available at SSRN 1694972, 2010. Also published in Risk, March 2011.

[2] Abraham Berman and Robert J. Plemmons. *Nonnegative Matrices in the Mathematical Sciences*, volume 9 of *Classics in Applied Mathematics*. SIAM, 1994. Reprint of the 1979 Academic Press edition.

[3] Hans Buehler. Expensive martingales. *Quantitative Finance*, 6(3):207–218, 2006.

[4] Hans Buehler. Volatility and dividends - volatility modelling with cash dividends and simple credit risk. `https://papers.ssrn.com/abstract=1141877`, 2008.

[5] Hans Buehler, Lukas Gonon, Josef Teichmann, and Ben Wood. Deep hedging. *Quantitative Finance*, 19:1271–1291, 2019. SSRN working paper version: doi:10.2139/ssrn.3120710 (2018).

[6] Hans Buehler, Blanka Horvath, Anastais Kratsios, Yannick Limmer, and Raeid Saqur. Sanos - smooth arbitrage-free non-parametric option surfaces. `https://arxiv.org/abs/2601.11209`, Februar 2026.

[7] Hans Buehler, Blanka Horvath, and Anastasis Kratsios. Dysanos generative dynamic smooth arbitrage-free non-parametric option surfaces, 2026.

[8] Hans Buehler and Evgenii Ryskin. Discrete local volatility for large time steps (short version). Technical report, SSRN, 2016. Short version. An extended version also exists: SSRN 2642630 (2015), doi:10.2139/ssrn.2642630.

[9] René Carmona and Sergey Nadtochiy. Local volatility dynamic models. *Finance and Stochastics*, 13:1–48, 2009.

[10] Samuel N. Cohen, Christoph Reisinger, and Sheng Wang. Arbitrage-free neural-sde market models, 2021. arXiv:2105.11053.

[11] Damir Filipović, Lane P. Hughston, and Andrea Macrina. Conditional density models for asset pricing. *International Journal of Theoretical and Applied Finance*, 15(1):1250002, 2012.

[12] Charles R. Johnson and Ronald L. Smith. Inverse M-matrices, II. *Linear Algebra and its Applications*, 435(5):953–983, 2011.

[13] Panayotis Mertikopoulos and Aris L. Moustakas. The emergence of rational behavior in the presence of stochastic perturbations. *The Annals of Applied Probability*, 20(4):1359–1388, 2010.

[14] P. E. Protter. *Stochastic Integration and Differential Equations*. Springer, 2 edition, 2005.

[15] Martin Schweizer and Johannes Wissel. Arbitrage-free market models for option prices: The multi-strike case. *Finance and Stochastics*, 12:469–505, 2008.

[16] Martin Schweizer and Johannes Wissel. Term structures of implied volatilities: Absence of arbitrage and existence results. *Mathematical Finance*, 18:77–114, 2008.

[17] Magnus Wiese, Lianjun Bai, Ben Wood, and Hans Buehler. Deep hedging: Learning to simulate equity option markets, 2019. NeurIPS 2019 Workshop on Robust AI in Financial Services.

[18] Magnus Wiese, Ben Wood, Alexandre Pachoud, Ralf Korn, Hans Buehler, Phillip Murray, and Lianjun Bai. Multi-asset spot and option market simulation, 2021. arXiv:2112.06823.

[19] Johannes Wissel. Arbitrage-free market models for option prices. NCCR FINRISK Working Paper 428, NCCR FINRISK, ETH Zürich, 2007.