Stochastic Smooth Discrete Local Volatility

Risk-Neutral Smooth Option Surfaces

Magnus Wiese* Hans Buehler* Blanka Horvath*,#

*Mathematical Institute, University of Oxford, #Oxford-Man Institute, University of Oxford

October 7, 2026

Abstract

We present a statically and dynamically arbitrage-free European option market model with smooth interpolation between option prices, based on the SANOS [BHK+26a] framework.

In spirit, it is a stocastic model for local volatility when the latter is used as a representation of a European option price surface, not as a pricing model.

We illustrate the strength of our approach on S&P Index Option data from January 2022 to August 2025. We show that the model's initial state fits within bid/ask for more than 95% of relevant market option prices.[1] We then train the model dynamics under $\mathbb{P}$ and $\mathbb{Q}$ and show that the model is indeed free of arbitrage while fitting the dynamics of the market well.

Contents

1 Introduction **2**

2 Stochastic Smooth Discrete Local Volatility **3**
- 2.1 Stochastic Discrete Local Volatility (SDLV) . . . 3
- 2.2 Risk-Neutral Dynamics . . . 4
  - 2.2.1 Future Densities . . . 4
  - 2.2.2 Front Density and Spot . . . 5
  - 2.2.3 Existence [TODO] . . . 7
- 2.3 Sampling . . . 7

3 Estimation **8**
- 3.1 Data and state representation . . . 9
- 3.2 End-to-end fitting . . . 9
- 3.3 DLV vol-of-vol estimation . . . 10

4 Absence of Dynamic Arbitrage **12**
- 4.1 Risk-Neutral Diagnostics . . . 13
- 4.2 Restricted Universe . . . 14
- 4.3 Full Strike Grid . . . 15

5 Conclusion **16**

[1] The model presented was trained for daily expiries up to 25 business days and strikes between $[\exp(-2.5\sigma_{atm}\sqrt{t}), \exp(+1.5\sigma_{atm}\sqrt{t})]$.

6 Thanks **16**

A Model and DLV factor visualizations **16**
A.1 Consistent Drift 16
A.2 PCA Summary 18

B Market-data and calibration pipeline **21**
B.1 Raw data and cleaning 21
B.2 Static preprocessing and quote selection 21
B.3 Daily calibration tensor 22
B.4 Optimization and retained artifacts 22

1 Introduction

This article presents an arbitrage-free market model for for smooth surfaces of European option prices. While not strictly the first of is kind, it is the first model with an industy-grade ability to fit the current market, allow for smooth interpolation of the entire option surface at each time, and provide an explicit construction of the continuous time risk-neutral drift for both (discrete local) volatilities and spot. It also has a natural estimation strategy for its parameter space.

The core architecture follows the SANOS approach [BHK$^+$26a] which models the stock price as the product of a time-and-state discrete process X and a smoothing process Y, in our case Black & Scholes. The former is used to encode the option market using a discrete form of local volatilities as code book. That does not mean that the model is a local volatility model – we merely use the implied local volatilities as a convenient representation of the option surface. Compared to using implied volatilities the advantage is that mere positivity of the local volatilities is sufficient to ensure the the option market is free of static arbitrage at all times. We have shown in [BHK$^+$26a] that this representation can be fitted to real life market data very efficiently using linear programming.

The key contribution in this paper is showing how to construct a dynamic model of SANOS option call prices which are not only statically but also dynamically free of arbitrage. We show the theoretical construction of the corresponding risk-neutral measure $\mathbb{Q}$ and discuss existence of the respective martingale processes. We then discuss implementation and estimate our model on S&P Index Option data from January 2022 to August 2025. We fit our daily model parameters to observed bid/ask prices and then proceed to estimate both $\mathbb{Q}$ and its drift under $\mathbb{P}$. We show that on this data set more than 90% of calibrated model prices fit within observed market bid/ask spreads even with the simplistic grid architecture we present here.

The model presented here is distinct from our own DYSANOS [BHK26b]. There, we focused on learning realistic market dynamics for smooth statically arbitrage-free SANOS surfaces under $\mathbb{P}$. We did also construct an equivalent risk-neutral measure $\mathbb{Q}$ numerically using the *near-martingale measure* approach introduced in [BMPW22]. While this method is model-free and can in theory be applied to any option market model under $\mathbb{P}$, it is numerically heavy and in practise restricted to making only a fairly limited number of the overall tradable instruments dynamically arbitrage-free.[2] In the current work we cover all options with an explicit construction of the risk-neutral drift, which is far more efficient than the numerical near-martingale measure approach.

[2] In [BHK26b] we used 49 instruments per simulated date.

Related Work

Our approach most closely related to the seminal but unpublished work [Wis07] which extends their earlier works [SW08a] and [SW08b]. From a theoretical perspective, our approach is a version of [CN09] where the authors derive related risk-neutral drift conditions for continuous time local volatility models, but without a tangible implementation strategy. Our results can be seen as a joint extension of the two papers, plus smoothing and market-fit from our SANOS [BHK$^+$26a].

In [CN12] the authors solve the problem of finding a dynamically arbitrage-free evolution of the full option surface via products of indepdent Levy processes. While structurally very elegant, option prices generated by products of independent increment processes do not typically fit well to equity index European option surfaces. Our approach also has the advantage of being numerically accessible. In particular, our model prices can be fit very efficiently within nearly all[3] of observed market bid/ask prices using the linear programming SANOS [BHK$^+$26a] calibration method.

There are a number of works on generative equity option market dynamics – not least our own 2019 work [WBWB19] and the multi-asset version [WWP$^+$21] – which simulate discrete strictly arbitrage-free option surfaces. However, most other approaches are not strictly free of arbitrage and do not have analytical corresponding risk-neutral dynamics. To address the absence of an analytical risk-neutral measure, we presented in [BMPW22] a model-free approach to learning such measure. However, that method is numerically challenging for large number of tradable options as discussed in [BHK26b].

For interest rate option markets the reference model in the same class is the BGM Model [BGM97].

2 Stochastic Smooth Discrete Local Volatility

In this section we introduce the core model architecture for our *Stochastic Smooth Discrete Local Volatility model*. We follow the core SANOS [BHK$^+$26a] idea of writing our equity process S as product

$$S_t = X_t Y_t \tag{1}$$

of a *Stochastic Discrete Local Volatility (SDLV)* process X which is used to fit to current market condition and whose state parameters will have consistent $\mathbb{P}$ and $\mathbb{Q}$ dynamics, and an independent[4] and rather trivial smoothing process Y whose dynamics under $\mathbb{P}$ and $\mathbb{Q}$ are the same.

The process Y will just be a Black & Scholes process with fixed volatility 5%. Given observed bid/ask market data it was shown in [BHK$^+$26a] how to fit a discrete local volatility surface in the presence of such process.

In the remainder of this section we focus on the SDLV process itself.

2.1 Stochastic Discrete Local Volatility (SDLV)

Our *Stochastic Discrete Local Volatility (SDLV)* risk-neutral model will be defined on a *fixed* expiry and cash strike grid. It will not use a floating grid in its option market dynamics.[5] We will use "t" to refer to model time. Denote initial spot by $X_0 > 0$.

[3] In below experiments, above 90%.

[4] Under both $\mathbb{P}$ and $\mathbb{Q}$.

[5] This is not to be confused to how the model parameters are estimated using historic market data. We explain estimation section 3.

For now let $0 = T_0 < T_1 < \cdots < T_M$ be the fixed model expiry dates, and assume that $K_0 < K_1 < \cdots < X_0 < \cdots < K_{N+1}$ are the fixed model cash strikes, where K^0 and K^{N+1} are boundary strikes which are far away from the inner strikes, c.f. [BHK$^+$26a].[6]

Assume that t is the current model time and let T_k be the front expiry such that $T_{k-1} \le t < T_k$. Our model is given at any time t by the current spot X_t and future densities q_t^j, where $q_t^{j,i}$ represents the risk-neutral density of the i-th cash strike at future expiry T_j.

Such discrete densities have been shown in [BR16] to admit an equivalent representation via (squared) *discrete local volatilities* $\Sigma_t^{j,1}, \ldots, \Sigma_t^{j,N}$ in the following form: there is a tridiagonal matrix Ω_j[7] for each $j > k$ such that

$$q_t^{j-1} = M_t^j q_t^j \ , \qquad M_t^j := I - \Omega_j S_t^j \ , \qquad S_t^j := \operatorname{diag}(0, \Sigma_t^{j,1}, \ldots, \Sigma_t^{j,N}, 0) \tag{2}$$

where M_t^j is an invertible tridiagonal M-matrix for every non-negative S_t^j. The key point is that q_t^j for $j > k$ are uniquely determined martingale densities for any given initial density q_t^k and the surface of non-negative (squared) discrete local vols.

2.2 Risk-Neutral Dynamics

Under a risk-neutral measure $\mathbb{Q}$ densities are martingales. We will now turn this property into conditions on dynamics of the (squared) discete local volatilities, and, jointly the front density and spot. Our approach follows in spirit that of [Wis07] in the sense that we rely on the martingale property of a discrete set of tradable derivative assets. However, we reframe it in a much more accessible coordinate system.

Assume that W is a d-dimensional Brownian motion. We designate W^1 as the spot factor: spot is always driven by dW_t^1.

2.2.1 Future Densities

Each density coordinate is a martingale. We therefore specify the front-density diffusion as

$$dq_t^k = E_t^{k;1}\, dW_t^1 \ , \qquad E_t^{k;1} := \operatorname{diag}(q_t^k)\eta_t^{k;1} \ , \qquad E_t^{k;f} := 0 \quad (f = 2, \ldots, d),$$

where $E_t^{k;1} \in \mathbb{R}^{N+2}$ must satisfy $\mathbf{1}^\top E_t^{k;1} = 0$, so that q_t^k retains unit mass. The remaining Brownian factors do not directly move the front density. They drive the stochastic DLV surface, while $\xi_t^{j,\cdot,1}$ determines Future density volatilities η_t^j for $j > k$ will be obtained recursively.

For future expiries T_j with $j > k$, assume the discrete local volatity squared is given as

$$\frac{d\Sigma_t^j}{\Sigma_t^j} = \mu_t^j\, dt + \xi_t^j\, dW_t,$$

[6] In our experiments the inner strikes are given as range of strikes in from $K^1 = \exp(-2.5\sigma_{atm}(T_M)\sqrt{T_M})$ to $K^N = \exp(+1.5\sigma_{atm}(T_M)\sqrt{T_M})$ with boundaries set to $K^0 = K^1/2$ and $K^{N+1} = 3K^N/2$; see section 4 for specifics.

[7] The matrix $\Omega_j \in \mathbb{R}^{(N+2)\times(N+2)}$ is given as follows: let $\gamma_t^{\pm i} := \pm(K^i)^2(T_j - T_{j-1})/(K^{i+1} - K^{i-1})/(K^{i\pm1} - K^i)$ for $i = 1, \ldots, N$. Then Ω_j is the tridiagonal matrix with entries

$$(\Omega_j)_{i,i-1} = \gamma_t^{-i}, \quad (\Omega_j)_{i,i} = -(\gamma_t^{-i} + \gamma_t^{+i}), \quad (\Omega_j)_{i,i+1} = \gamma_t^{+i}, \quad i = 1, \ldots, N.$$

c.f. [BHK$^+$26a].

where $\xi_t^j \in \mathbb{R}^{N\times d}$ is given and $\mu_t^j \in \mathbb{R}^N$ is the unknown risk-neutral drift. We write this as

$$dS_t^j = A_t^j\,dt + \sum_{f=1}^{d} B_t^{j;f}\,dW_t^f,$$

with $A_t^j := \operatorname{diag}(0, \Sigma_t^{j,1}\mu_t^{j,1}, \ldots, \Sigma_t^{j,N}\mu_t^{j,N}, 0)$ and $B_t^{j;f} := \operatorname{diag}(0, \Sigma_t^{j,1}\xi_t^{j,1;f}, \ldots, \Sigma_t^{j,N}\xi_t^{j,N;f}, 0)$. With the above

$$dM_t^j = -\Omega_j\,dS_t^j = -\Omega_j A_t^j\,dt - \sum_{f=1}^{d}\Omega_j B_t^{j;f}\,dW_t^f.$$

Setting $dq_t^j = \sum_f E_t^{j;f}\,dW_t^f$ for an unknown $E_t^{j;f}$ yields

$$\begin{aligned} dq_t^{j-1} &= M_t^j\,dq_t^j + (dM_t^j)q_t^j + d[M^j, q^j]_t && (3)\\ &= \sum_{f=1}^{d}\left(M_t^j E_t^{j;f} - \Omega_j B_t^{j;f} q_t^j\right)dW_t^f - \Omega_j\left(A_t^j q_t^j + \sum_{f=1}^{d} B_t^{j;f}E_t^{j;f}\right)dt\,. && (4)\end{aligned}$$

Matching the diffusion coefficients gives

$$E_t^{j;f} = (M_t^j)^{-1}\left(E_t^{j-1;f} + \Omega_j B_t^{j;f} q_t^j\right) \in \mathbb{R}^{N+2} \quad \text{and} \quad \eta_t^{j,i;f} = \frac{E_t^{j,i;f}}{q_t^{j,i}}\,. \tag{5}$$

The drift must vanish, hence

$$\Omega_j\left(A_t^j q_t^j + \sum_{f=1}^{d} B_t^{j;f}E_t^{j;f}\right) = 0.$$

The restriction of Ω_j to vectors supported on the interior nodes is injective. Therefore the bracket vanishes node by node, giving

$$\mu_t^{j,i} = -\sum_{f=1}^{d}\xi_t^{j,i;f}\frac{E_t^{j,i;f}}{q_t^{j,i}} = -\sum_{f=1}^{d}\xi_t^{j,i;f}\eta_t^{j,i;f}, \qquad i = 1,\ldots,N. \tag{6}$$

Thus the prescribed front-density volatility and the prescribed DLV volatilities determine all future density volatilities and the risk-neutral DLV drifts by forward induction over the expiries.

2.2.2 Front Density and Spot

The remaining task is to specify both spot dynamics consistent with the front density and the respective front density volatilities.

Case $t = T_k - dt$

Assume we are at $t = T_k - dt$ and observe the front density q_t^k. Current spot is given as $S_t = K^\top q_t^k$. In order to sample a spot consistent with the discrete set of strikes, we can either draw a discrete sample or fit some other distribution exactly to the observed option prices (fitting to densities or probabilities will not work in general).

The simplest method is to draw spot from this discrete distribution using the spot factor W^1. Let F_t^k denote the conditional CDF which assigns probability $q_t^{k,i}$ to strike K_i, and let Φ denote the standard normal CDF. Then

$$S_{T_k} = (F_t^k)^{-1}(\Phi(Z))\,, \qquad Z := \frac{W_{T_k}^1 - W_t^1}{\sqrt{dt}} \sim N(0,1).$$

Correlation between spot and a future DLV coordinate is specified by that coordinate's loading $\xi_t^{j,i;1}$ on W^1.

Conditional on this definition of terminal spot, the front-density diffusion is exact if the Brownian increment is revealed progressively: define

$$Q_i := \sum_{\ell=0}^{i} q_t^{k,\ell}, \qquad a_i := \Phi^{-1}(Q_i), \qquad a_{-1} := -\infty, \quad a_{N+1} := +\infty,$$

and let ϕ be the standard normal density. For $t \le s < T_k$, set

$$z_i(s) := \frac{\sqrt{dt}\,a_i - (W_s^1 - W_t^1)}{\sqrt{T_k - s}}.$$

The time-s front density is the conditional probability

$$q_s^{k,i} := \mathbb{P}(S_{T_k} = K_i \mid \mathcal{F}_s) = \Phi(z_i(s)) - \Phi(z_{i-1}(s))\,. \tag{7}$$

It is therefore a martingale. Applying Itô's formula gives the exact diffusion

$$dq_s^{k,i} = E_s^{k,i;1}\,dW_s^1, \qquad E_s^{k,i;1} = \frac{\phi(z_{i-1}(s)) - \phi(z_i(s))}{\sqrt{T_k - s}}, \qquad t \le s < T_k, \tag{8}$$

with zero drift. Since $z_i(t) = a_i$, its initial value is

$$E_t^{k,i;1} = \frac{\phi(a_{i-1}) - \phi(a_i)}{\sqrt{dt}} \quad \text{hence} \quad \eta_t^{k,i;1} = \frac{\phi(a_{i-1}) - \phi(a_i)}{q_t^{k,i}\sqrt{dt}} \quad \text{for } i = 0, \dots, N+1. \tag{9}$$

while $E_t^{k,i;f} = \eta_t^{k,i;f} = 0$ for $f = 2, \dots, d$. More generally, $\eta_s^{k,i;1} = E_s^{k,i;1}/q_s^{k,i}$ and the other factor loadings remain zero. Consequently,

$$S_s = K^\top q_s^k = \mathbb{E}[S_{T_k} \mid \mathcal{F}_s]\,, \qquad dS_s = K^\top E_s^{k;1}\,dW_s^1,$$

so spot and the front density are driven by the same Brownian factor. At expiry, q_s^k converges to the point mass associated with the sampled strike. The displayed coefficient $E_t^{k,i;1}$ is only the initial value of the exact, state- and time-dependent coefficient $E_s^{k,i;1}$; holding it constant over the whole interval would instead be an Euler approximation. Likewise, sampling S_{T_k} all at once without progressively revealing W^1 produces a direct jump to the terminal point mass rather than a continuous density diffusion.

Case $t \ll T_k$

The same construction gives a natural exact transition over any step which does not cross the front expiry. Let

$$h := T_k - t, \qquad 0 < \delta \le h, \qquad \Delta W_t^1 := W_{t+\delta}^1 - W_t^1.$$

Starting from q_t^k, define Q_i and a_i as above. After drawing $\Delta W_t^1 \sim N(0, \delta)$, set

$$q_{t+\delta}^{k,i} = \Phi\left(\frac{\sqrt{h}\, a_i - \Delta W_t^1}{\sqrt{h-\delta}}\right) - \Phi\left(\frac{\sqrt{h}\, a_{i-1} - \Delta W_t^1}{\sqrt{h-\delta}}\right), \qquad \delta < h. \tag{10}$$

The corresponding spot is not sampled independently but is defined by

$$S_{t+\delta} := K^\top q_{t+\delta}^k = \mathbb{E}[S_{T_k} \mid \mathcal{F}_{t+\delta}]\,. \tag{11}$$

Consequently

$$\mathbb{E}_t[q_{t+\delta}^k] = q_t^k, \qquad \mathbb{E}_t[S_{t+\delta}] = S_t$$

exactly. In general $S_{t+\delta}$ is not a grid strike; only at T_k does the density collapse to a unit vector and spot become the corresponding strike.

This update is recursively consistent. Indeed, the Gaussian boundaries calculated from the cumulative probabilities of $q_{t+\delta}^k$ are precisely the residual boundaries appearing in (10). The construction may therefore be restarted from $q_{t+\delta}^k$ without changing the joint law. It is a simplex-valued information bridge: conditional on the exact terminal Gaussian variable, W^1 is a Brownian bridge; conditional only on the terminal strike, it is Brownian motion conditioned to terminate in an interval, or equivalently a mixture of Brownian bridges.

2.2.3 Existence [TODO]

We have just described a set of risk-neutral model dynamics for a discrete process X_t whose future densities themselves are martingales driven by stochastic discrete local volatilities. Our construction is concise and lends itself to efficient numerical implementation and model parameter estimation.

[TODO] Existience

2.3 Sampling

In this section we discuss the sampling procedure for the discrete local volatility model. The final model is a product of this process and a simple Black & Scholes process with, say, volatility of 5%. It serves as a SANOS [BHK+26a] smoother and ensures that the overall option price surface is smooth in strike.

On a time grid $t_n < t_{n+1}$ with $\delta_n = t_{n+1} - t_n$, one may proceed as follows.

First draw one Brownian vector $\Delta W_n \in \mathbb{R}^d$, using its first coordinate both in (10) and in every DLV coordinate with a non-zero loading $\xi^{j,i;1}$. This preserves the specified leverage. Update the front density by (10), and calculate spot from (11). This part of the update has no time-step bias, preserves positivity and unit mass, and is exactly martingale preserving.

For the DLV squared it is useful to work with $X_t^{j,i} := \log \Sigma_t^{j,i}$. Its Itô dynamics are

$$dX_t^{j,i} = \left(\mu_t^{j,i} - \frac{1}{2}\sum_{f=1}^{d} (\xi_t^{j,i;f})^2\right) dt + \sum_{f=1}^{d} \xi_t^{j,i;f}\, dW_t^f\,.$$

Even the frozen-coefficient exponential update

$$\Sigma_{n+1}^{j,i} = \Sigma_n^{j,i} \exp\left\{ \left(\mu_n^{j,i} - \frac{1}{2} \sum_f (\xi_n^{j,i;f})^2 \right) \delta_n + \sum_f \xi_n^{j,i;f} \Delta W_n^f \right\} \tag{12}$$

is preferable to Euler integration of Σ: it cannot create a negative local variance. For state-dependent coefficients, a genuine Itô weak-order-two stochastic Runge–Kutta method in the log variables reduces the remaining time-step bias.

After updating Σ^j reconstruct the densities expiry by expiry by solving the tridiagonal systems

$$q_{n+1}^j = \left(I - \Omega_j S_{n+1}^j \right)^{-1} q_{n+1}^{j-1}, \qquad j = k+1, \ldots, M. \tag{13}$$

This enforces the DLV relation at every simulated state. Because the matrix is an M-matrix for non-negative Σ, the solve preserves positivity.

The drift condition (6) removes the infinitesimal drift of every density coordinate. It does not, by itself, make a generic finite-step discretisation exactly unbiased, since the map $\Sigma^j \mapsto (M^j)^{-1} q^{j-1}$ is nonlinear. Thus (12), or any ordinary higher-order SDE scheme, retains a time-discretisation bias which should vanish under step refinement.

Our experiments in 4 use this sampling scheme with antithetically paired Brownian increments.

Remark 2.1 (Alternative discrete-branch construction). If exact martingality under a chosen finite-step simulation law is required, it can instead be imposed directly. Choose symmetric discrete shock branches $\Delta W_{n,b}$ with probabilities w_b and construct the front-density branches so that

$$\sum_b w_b q_{n+1,b}^k = q_n^k \, .$$

Parameterise each positive DLV update with an unknown drift correction $c_n^j \in \mathbb{R}^N$. For example, replace the log drift in (12) by $c_n^{j,i}$. For each expiry in turn, solve the nonlinear system

$$\sum_b w_b q_{n+1,b}^j(c_n^j) = q_n^j, \qquad q_{n+1,b}^j(c_n^j) = \left(I - \Omega_j S_{n+1,b}^j(c_n^j) \right)^{-1} q_{n+1,b}^{j-1}. \tag{14}$$

Only N coordinates are independent because mass and first moment are already preserved, matching the N drift corrections. Initialising the nonlinear solver with the continuous-time drift (6) gives a stable finite-step correction. Equation (14) makes the densities exactly martingales under the discrete branch law; symmetric or antithetic branches reduce variance but do not replace this condition.

3 Estimation

We briefly describe out estimation strategy. There are two components: fitting daily the "implied" discrete densities and therefore discete local volatilities to observed bid/ask spreads, and then an estimation of the vol of those discrete local volatilities, and their statistical drift.

The former represents today's fit of the model to the market, while the latter captures the evolution of the model parameters over time.

The summary here is brief; refer to DYSANOS [BHK26b] for a comprehensive discussion on parameter mapping and estimation considerations.

3.1 Data and state representation

We estimate a family of models which has 25 expiries at one-business-day intervals. The model is aimed as S&P 500 index options whose main trading activity occurs in the short-dated options segment.

Each daily state has 25 model expiries at one-business-day intervals and 200 interior strike nodes, together with two artificial boundary nodes. The 200 inner nodes are generated from fixed standardized log-moneyness templates on $[-2.5, 1.5]$.[8] They therefore share one strike support across the 25 expiries on a given date, while that support floats between dates with the observed ATM term structure. A fixed 5% annual Black volatility supplies the independent SANOS smoothing factor.

We use European S&P 500 index-option quotes via WRDS[9] from OptionMetrics IvyDB on 891 valid market dates from 3 January 2022 to 29 August 2025. After the cleaning and eligibility pipeline described in appendix B, the retained sample contains 2,192,458 original call and put bid–ask intervals with at most 25 business days to expiry. Prices and strikes are forward-normalized, while the contemporaneous spot is retained separately.

3.2 End-to-end fitting

Let $x_d \in \mathbb{R}^{25\times 200}$ denote the unconstrained state on date d. Its DLV-squared coordinates are decoded as

$$\Sigma_d^{j,i} = \min\{9, 10^{-4} + \exp(x_d^{j,i})\}.$$

The SANOS recursion then produces all densities and quote prices. Every retained market quote is priced at its own strike and expiry; quotes between two model expiries use the corresponding linear time interpolation. The fit is therefore performed directly against the original market bid and ask, not against an intermediate linear-programming state.

The *direct* specification fits all 5,000 coordinates of x_d independently for every date. The *affine* specifications impose the shared representation

$$x_d = \bar{x} + \sum_{a=1}^{r} z_d^a u_a, \qquad r \in \{100, 200, 400\}, \tag{15}$$

and jointly estimate the mean $\bar{x}$, loadings u_a, and daily codes z_d. See [BHK26b] for motivation and further results.

For model price C, quoted interval $[b, a]$, midpoint $m = (a+b)/2$, and half-spread $h = (a-b)/2$, define the bid–ask *excess*

$$e(C; b, a) = \frac{\max(C - a, 0) + \max(b - C, 0)}{h}.$$

[8] The expiry grid is fixed at $T_j = j/252$, $j = 1, \ldots, 25$. The strike construction uses 100 fixed normalized-moneyness coordinates for the one-day ATM scale and 100 for the 25-day ATM scale, both on $[-2.5, 1.5]$. On each calibration date these coordinates are mapped to log-relative-strike candidates using that date's two regularized ATM total-volatility scales, merged and sorted, and given a minimum gap equal to 5% of the average log-strike spacing. Exponentiation gives one 200-node relative-strike grid shared by all 25 expiries; the boundary strikes are one half of the lowest and one and a half times the highest inner strike. Thus the expiry and normalized-moneyness templates are fixed, and the resulting cash strikes are fixed along each simulated path. For historical calibration the relative-strike grid floats between dates because the ATM scales change.

[9] `https://wrds-www.wharton.upenn.edu/`

Writing n_d for the number of quotes on date d, the quote component of the training objective is

$$\mathcal{L} = \frac{1}{891} \sum_d \frac{1}{n_d} \sum_{\ell=1}^{n_d} \left[e(C_{d\ell}; b_{d\ell}, a_{d\ell}) + 0.05 \frac{|C_{d\ell} - m_{d\ell}|}{h_{d\ell}} \right], \tag{16}$$

plus a soft barrier near the upper DLV-squared bound, c.f. [BHK26b]. Thus a price within its quoted interval has zero excess, while the small midpoint term identifies a stable point inside the spread. We use the exact all-quote gradient, accumulated in memory-bounded quote blocks, and Adam for up to 1,000 epochs.

Table 1 reports statistics pooled over all 2,192,458 retained quotes. Conditional excess is the mean of e only over quotes outside their spread; midpoint MAE is also measured in half-spreads.

Representation	State dim.	Inside bid–ask	Mean excess	Conditional excess	Midpoint MAE
Direct 25-by-200	5,000	97.32%	0.035	1.288	0.237
Affine-100	100	89.68%	0.083	0.808	0.471
Affine-200	200	95.02%	0.045	0.900	0.349
Affine-400	400	96.01%	0.040	1.009	0.317

Table 1: End-to-end fit to the retained market bid–ask intervals. State dimension is 5,000 for the unrestricted daily state and the latent rank for an affine state.

The direct representation gives the highest bid–ask coverage and the lowest midpoint error and is therefore used for the stochastic-DLV experiments below. Affine-200 already exceeds 95% bid–ask coverage; increasing the rank to 400 gives a smaller further improvement. Conditional excess is computed only on the changing set of failed quotes and should not be read as a monotone ranking across representations.

3.3 DLV vol-of-vol estimation

The floating historical supports cannot be differenced node by node without first aligning contracts. We use 710 adjacent market-session pairs ending no later than 31 December 2024. Write

$$L_d^j(x) := \log \Sigma_d^j(x), \qquad R_d := \frac{S_{d+1}}{S_d},$$

and let $x_{d,i}$ be the ith relative strike on the origin-date support. In the zero-carry normalization used for the dynamics, holding the corresponding cash strike fixed moves its relative strike to $x_{d,i}/R_d$ on the next date. Holding calendar expiry fixed simultaneously shifts the remaining-expiry index down by one. The aligned observation is therefore

$$\Delta L_d^{j,i} = L_{d+1}^j \left(\frac{x_{d,i}}{R_d} \right) - L_d^{j+1}(x_{d,i}), \qquad j = 1, \ldots, 24, \quad i = 1, \ldots, 200. \tag{17}$$

The endpoint value is interpolated on date $d+1$'s own floating support. Thus cash strike and calendar expiry are fixed within each one-day pair, while the origin chart is allowed to float across different pairs.

For each expiry–strike coordinate we fit the physical one-step regression

$$L_{d+1}^j \left(\frac{x_{d,i}}{R_d} \right) = a^{j,i} + b^{j,i} L_d^{j+1}(x_{d,i}) + \varepsilon_d^{j,i}. \tag{18}$$

The centered residual surfaces are smoothed by a Gaussian filter with widths one expiry node and two strike nodes. Let

$$Z_d = \Phi^{-1}\big(F_d^1(R_d)\big), \qquad \widetilde{Z}_d = \frac{Z_d - \overline{Z}}{s_Z},$$

where F_d^1 is the fitted front distribution including the independent 5% Black factor. The daily loading of the spot factor is

$$\ell_1^{j,i} = \frac{1}{n}\sum_d \widetilde{\varepsilon}_d^{j,i}\,\widetilde{Z}_d.$$

We remove this component and apply PCA to $\widetilde{\varepsilon}_d - \widetilde{Z}_d\ell_1$; its leading eigenvectors give the remaining daily loading surfaces $\ell_2, \ell_3, \ldots$. The annualized DLV vol-of-vol used in the SDE is

$$\xi^{j,i;f} = \sqrt{252}\,\ell_f^{j,i}. \tag{19}$$

The fitted spot scores have mean 0.0290 and standard deviation 1.0026. Eight factors retain 72.88% of the smoothed innovation variance and 16 retain 84.40%; 26 factors are required to exceed 90%, reaching 90.25%. The rank-16 system is nested, with its first eight loading surfaces identical to the rank-8 system. Appendix A.2 shows the baseline implied-volatility surface, the cumulative factor power, and the implied- volatility response to each retained factor. On the held-out 2025 sample the full 25-by-200 representation has residual root mean square 0.256, compared with 0.423 for the affine-x 200 representation.

This is a functional estimate in remaining-expiry and normalized-moneyness coordinates, not an estimate on one global historical cash-strike grid. When a simulation is initialized, these loading surfaces are attached nodewise to that state's fixed cash-strike grid. Pooling dates therefore assumes that the loading functions are stable in this moving coordinate chart.

The empirical physical log-DLV drift implied by regression (18) is

$$\widehat{\mu}_{d,k}^{\mathbb{P},\mathrm{emp}} = 252\,[a_k + (b_k - 1)L_{d,k}], \tag{20}$$

where k denotes an aligned expiry–strike coordinate. At the same historical state we compute $\mu_d^{\mathbb{Q}}$ from the martingale drift restriction in (6). Under the convention $\mathrm{d}W^{\mathbb{Q}} = \mathrm{d}W^{\mathbb{P}} + \lambda\,\mathrm{d}t$, a physical drift compatible with the fitted diffusion must have the form

$$\mu_d^{\mathbb{P},\mathrm{comp}} = \mu_d^{\mathbb{Q}} + \xi\lambda_d. \tag{21}$$

We estimate the state-dependent market price of risk by

$$\widehat{\lambda}_d = \arg\min_{\lambda\in\mathbb{R}^r} \sum_k w_{d,k}\left[\widehat{\mu}_{d,k}^{\mathbb{P},\mathrm{emp}} - \mu_{d,k}^{\mathbb{Q}} - (\xi\lambda)_k\right]^2. \tag{22}$$

Here $w_{d,k}$ is proportional to the interior risk-neutral density, normalized so that each expiry has equal total weight. We also estimate one constant $\widehat{\bar{\lambda}}$ by pooling the same objective over all dates; this is the measure change used in the simulations. The independent 5% Black factor is assigned zero risk premium and therefore has the same law under $\mathbb{P}$ and $\mathbb{Q}$.

For comparison with a model-free time-series estimate, define the unconditional coordinate drift

$$\widehat{\bar{\mu}}_k^{\mathbb{P},\mathrm{unc}} = \frac{252}{n}\sum_{d=1}^{n}\Delta L_{d,k}. \tag{23}$$

Because (18) contains an intercept, this is also the sample mean of (20). We measure the distance of a compatible drift from this unconditional estimate by

$$D = \left[\frac{\sum_k \bar{w}_k \left(\overline{\mu^{\mathbb{P},\mathrm{comp}}}_k - \widehat{\bar{\mu}}_k^{\mathbb{P},\mathrm{unc}} \right)^2}{\sum_k \bar{w}_k} \right]^{1/2}, \qquad \bar{w}_k = \frac{1}{n} \sum_{d=1}^{n} w_{d,k}. \tag{24}$$

Calibration quantity	Rank 8	Rank 16
PCA innovation variance retained	72.88%	84.40%
$\mathbb{Q}$-drift RMS	13.58	13.77
Empirical $\mathbb{P}-\mathbb{Q}$ target RMS	39.41 (38.95, 39.88)	39.29 (38.82, 39.75)
Statewise compatible residual RMS	24.23 (24.01, 24.44)	20.81 (20.63, 20.99)
Statewise target energy captured	62.20% (61.64, 62.75)	71.95% (71.52, 72.37)
Constant-λ residual RMS	29.97 (29.58, 30.37)	28.54 (28.14, 28.94)
Constant-λ target energy captured	42.15% (40.94, 43.32)	47.24% (46.10, 48.37)
D to unconditional drift, statewise λ_d	16.33 (15.75, 16.90)	13.57 (13.14, 14.01)
D to unconditional drift, constant $\bar{\lambda}$	16.43 (15.80, 17.05)	13.73 (13.21, 14.26)

Table 2: Vol-of-vol coverage and physical-drift projection on the 710 one-session training pairs. Parentheses give 95% intervals. RMS quantities are annualized log-DLV-squared drifts. The target and residual intervals use 20,000 pair-bootstrap resamples; the intervals for D use the pair-level delta method. They condition on the fitted AR, PCA, and risk-neutral drift, so they quantify time-series sampling error rather than full re-estimation uncertainty.

The statewise fit is a diagnostic upper bound on what the retained diffusion span can reproduce. The constant-λ rows describe the more restrictive measure change actually used below. Rank 16 materially improves both fits, but neither retained factor system reproduces the empirical physical drift exactly. In particular, the confidence intervals in the last two rows do not include zero; the compatible drift remains measurably different from the raw unconditional estimate.

Appendix A.1 visualizes the effect of the drift change with four matched Model D paths.

4 Absence of Dynamic Arbitrage

In this section we assess whether the model is indeed free of dynamic arbitrage. We are not aware of an accessible benchmark for dynamic arbitrage in the literature, so the methodology described here is a contribution in its own right.

We compare four model specifications:

Model	DLV factors	Steps/day
A	8	1
B	16	1
C	8	8
D	16	8

Every statistical audit uses 100,000 paths arranged as 50,000 independent antithetic pairs. The

horizon is 20 business days. For a live instrument i, the one-day return used in both diagnostics is

$$R^i_{p,n} = \frac{V^i_{p,n+1} - V^i_{p,n}}{Z_{p,n}}, \tag{25}$$

where Z is spot and V^i is either spot or an option price. Division by the pre-step spot is predictable and therefore preserves the zero-conditional-mean property under $\mathbb{Q}$.

4.1 Risk-Neutral Diagnostics

Numerically, it is clear that as the number of instruments increases on a fixed set of paths, the likelihhood of a risk-free return of a linear combination of instruments simply due to sampling error increases.

We will therefore test absence of dynamic arbitrage on a reduced trading universe which consists of spot and 44 options with cash strikes fixed at the initial date: eleven OTM calls or puts at each of 5, 10, 15, and 20 business days. Thus 45 instruments are live initially, giving 20 spot and 550 option day–instrument coordinates over the 20-day horizon.

The resulting test will only confirm absence of dynamic arbitrage within this restricted set of instruments, and only within the class of strategies considered. To illustrate the limitations of the former, we apply the same test to an enlarged trading universe in section 4.3.

Unconditional diagnostic. For every day–instrument coordinate (n, i), first average the two members of each antithetic pair and then form

$$z_{n,i} = \frac{\overline{R}_{n,i}}{\widehat{\mathrm{se}}(\overline{R}_{n,i})}. \tag{26}$$

The tables report the maximum and 99th percentile of $|z_{n,i}|$ and the number of coordinates beyond fixed thresholds. This checks the unconditional implication $\mathbb{E}^{\mathbb{Q}}[R^i_{n+1}] = 0$ separately at every day and instrument. It does not test $\mathbb{E}^{\mathbb{Q}}[R^i_{n+1} \mid \mathcal{F}_n] = 0$ at every simulated state. Under $\mathbb{P}$ the same zero-mean statistics describe fitted physical drift; they are not no-arbitrage violations. For one standard-normal coordinate the 99th percentile of $|z|$ is 2.576. Coordinate dependence and the simultaneous screen mean that the reported counts are diagnostics, not family-wise p-values. The 1.96 and 2.576 thresholds apply to each coordinate-level z-score, not directly to the maximum or empirical quantiles across all coordinates.

Conditional diagnostic. We split complete antithetic pairs into two equal path sets A and B, and write $\mathcal{I}_A$ for all live path–time observations from paths in A. The state $X_{p,n}$ contains the cumulative Brownian-factor histories, log spot, the live log option-price-to-spot ratios, calendar time, and time to each expiry. On set A we fit the uncapped affine-plus-deep action

$$a_\theta(X_n) = D_n \left[b + L\widetilde{X}_n + g_\theta(\widetilde{X}_n) \right], \tag{27}$$

where D_n masks expired instruments, $\widetilde{X}$ is standardized using set A only, and g_θ is a two-layer width-32 SiLU network. The training objective is

$$\widehat{L}_A(\theta) = \log \left[\frac{1}{|\mathcal{I}_A|} \sum_{(p,n)\in\mathcal{I}_A} \exp\left(-a_\theta(X_{p,n})^\top R_{p,n} + \frac{1}{2}\|a_\theta(X_{p,n})\|^2 \right) \right]. \tag{28}$$

For conditionally Gaussian returns the population action is approximately[10] $(I+\mathrm{Cov}_t(R))^{-1}\mathbb{E}_t[R]$, so the fit targets a regularized covariance-adjusted conditional mean.

The fitted action is frozen and evaluated on the untouched set B. Its score on path p is the average daily captured return

$$d_p = \frac{1}{20}\sum_{n=0}^{19} a_\theta(X_{p,n})^\top R_{p,n}. \tag{29}$$

The factor $1/20$ only fixes the unit; removing it multiplies both the estimate and its standard error by 20. We average d_p within each antithetic pair and report its mean divided by its pair-based standard error as the *action statistic*, denoted T_{act}. We then exchange A and B and combine the two out-of-fold scores.

Under the conditional zero mean null hypothesis every frozen predictable action has zero expected score, and T_{act} is approximately standard normal. Thus $|T_{\mathrm{act}}| < 1.96$ is consistent with the null at the 95% level and $|T_{\mathrm{act}}| < 2.576$ at the 99% level. Accordingly, an action statistic is shown in black for $|T_{\mathrm{act}}| \le 1.96$, orange for $1.96 < |T_{\mathrm{act}}| \le 2.576$, and red for $|T_{\mathrm{act}}| > 2.576$. These are per-test confidence thresholds. This is a conditional test relative to the state and strategy class in (27), rather than a proof over all measurable strategies.

For broad strike universes, raw z-scores can be dominated by tail calls whose prices and Monte Carlo standard errors are both microscopic. To distinguish statistical from economic size, write $\widehat{m}_i$ and $\widehat{\mathrm{SE}}_i$ for the pair-based mean and standard error and define the materiality score

$$s_i^{(\varepsilon)} = \frac{\widehat{m}_i}{\sqrt{\widehat{\mathrm{SE}}_i^2 + \varepsilon^2}}, \qquad \varepsilon = 10^{-8}. \tag{30}$$

This is not a z-statistic: ε places an absolute floor, in spot units, beneath economically negligible sampling errors. We use this additional diagnostic only for the broad-grid extensions below.

4.2 Restricted Universe

Table 3 shows the $\mathbb{Q}$ results for the restricted trading universe. Orange marks a moderate departure from the Gaussian benchmark and red would mark a five-standard-error violation in the coordinate-wise screen. The separate action-statistic rule is given above; all four T_{act} entries are black.

Model	N	Max. $\lvert z\rvert$	99% $\lvert z\rvert$	$N_{2.576}$	N_5	T_{act}
		Unconditional				
A: rank 8, one step	570	2.716	2.464	3	0	1.401
B: rank 16, one step	570	3.663	3.618	10	0	1.040
C: rank 8, eight steps	570	2.793	2.373	4	0	1.753
D: rank 16, eight steps	570	2.996	2.967	22	0	0.040

Table 3: Restricted-universe, 100,000-path $\mathbb{Q}$ diagnostics. The first five numeric columns are coordinate-wise unconditional statistics; the last is the two-fold out-of-fold conditional action statistic. Here N_c is the number of coordinates with $|z| > c$.

No restricted coordinate exceeds five standard errors, and no conditional action statistic exceeds 1.753 in absolute value. Model D has the richest factor representation, the finer time

[10] On the filtration generated by the tangent space of the optimal action.

discretisation, and the conditional statistic closest to zero. Table 4 applies the same diagnostics to paths generated with the fitted physical drift. The entries are deliberately not colour coded: nonzero physical expected returns are intended evidence of the measure change, not arbitrage failures.

Model	N	Max. $\lvert z\rvert$	99% $\lvert z\rvert$	$N_{2.576}$	N_5	T_{act}
		Unconditional				
A: rank 8, one step	570	62.138	59.780	570	570	90.038
B: rank 16, one step	570	342.470	328.811	570	570	646.645
C: rank 8, eight steps	570	61.002	58.398	570	570	89.205
D: rank 16, eight steps	570	336.819	325.038	570	570	624.560

Table 4: Restricted-universe, 100,000-path $\mathbb{P}$ diagnostics, reported against zero to display the predictable drift created by the fitted measure change.

All 570 coordinates exceed five standard errors under every physical-measure specification, and the out-of-fold action statistics range from 89.205 to 646.645. Thus the same diagnostics that remain small under $\mathbb{Q}$ recover both unconditional and state-predictable drift after the fitted change to $\mathbb{P}$.

4.3 Full Strike Grid

The preceding experiment deliberately trades few options. We next include spot and all 200 interior fixed-strike calls at each of 5, 10, 15, and 20 business days. There are 801 instruments at the initial date and 10,020 live day–instrument coordinates. We therefore supplement the raw statistics with the materiality score in (30).

Model	Max. $\lvert z\rvert$	99% $\lvert z\rvert$	N_5	Max. $\lvert s^{(\varepsilon)}\rvert$	99% $\lvert s^{(\varepsilon)}\rvert$	T_{act}
			Unconditional			
A: rank 8, one step	19.181	2.357	13	2.738	2.283	−0.218
B: rank 16, one step	13.360	2.243	26	3.240	2.133	0.051
C: rank 8, eight steps	13.318	2.631	13	4.481	2.578	2.175
D: rank 16, eight steps	9.043	2.761	19	4.784	2.759	−0.562

Table 5: Full-grid, 100,000-path $\mathbb{Q}$ diagnostics. Red entries are raw five-standard-error tail violations; orange entries are moderate raw or materiality-score deviations. No materiality score exceeds five.

Thus the clean restricted-universe result does depend on excluding the most extreme strikes. Every full-grid model has raw unconditional tail violations, although the conditional action statistics remain small and no materiality-regularized score exceeds five. For example, the worst Model A coordinate is the highest-strike five-day call: its initial value is 2.31×10^{-9} spot units, its mean one-day increment is 4.48×10^{-10}, and only 0.056 exercises are expected among 100,000 paths. The raw $z = 19.18$ identifies a finite-step tail defect, while (30) records its negligible absolute scale.

Interpretation. The restricted book produces acceptable $\mathbb{Q}$ diagnostics for every configuration, with Model D closest to zero in the conditional action test. The full option grid gives the more demanding conclusion: all four finite-step schemes retain statistically detectable defects in tail calls. These are numerical defects of the discretisation, not of the continuous-time drift identity; see Section 2.3.

5 Conclusion

We have given an explicit continuous-time construction of a risk-neutral measure for the SANOS discrete local volatility representation of European option markets, extending the static SANOS surface of [BHK$^+$26a] to a dynamically arbitrage-free model. The resulting SDLV process admits a closed-form risk-neutral drift for every discrete local volatility coordinate and an exact front-density/spot diffusion, in place of the numerical near-martingale measure used in [BHK26b].

Estimated on 891 trading days of S&P 500 index option data, the model's static fit covers more than 95% of quoted bid/ask intervals, and the restricted-universe $\mathbb{Q}$ diagnostics of Section 4 confirm the theoretical martingale property up to time-discretisation error. On the full option grid, finite-step simulation retains statistically detectable but economically negligible tail defects, consistent with the discretisation bias discussed in Section 2.3 rather than a failure of the continuous-time drift identity. We further estimate the physical dynamics and the implied market price of risk, recovering a stable low-rank factor structure for the DLV vol-of-vol.

Natural next steps include a finer or higher-order time-stepping scheme to reduce the residual tail bias documented in Section 4.3 (c.f. the discrete-branch correction of Remark 2.1), a formal existence and uniqueness argument for the risk-neutral SDE system of Section 2.2.1, and extending the estimation beyond short-dated S&P 500 index options to other asset classes and longer tenors.

6 Thanks

We want to thank the team at Wharton for their WRDS `https://wrds-www.wharton.upenn.edu` portal and Option Metrics `https://optionmetrics.com/united-states` for providing their IvyDB data via this service. Without this data our work would not have been possible.

A Model and DLV factor visualizations

A.1 Consistent Drift

Figure 1 illustrates the fitted measure change using four matched Model D paths. Solid and dashed paths of the same colour use identical Gaussian draws under $\mathbb{Q}$ and $\mathbb{P}$. Solid and dashed lines of the same colour use the same Gaussian draws under $\mathbb{Q}$ and $\mathbb{P}$. Besides spot, the illustration reports ATM one- and five-day implied volatility and the central log-strike skew

$$\mathrm{Skew}_\tau = \frac{\sigma_{\mathrm{imp}}(K = Z_t e^h, \tau) - \sigma_{\mathrm{imp}}(K = Z_t e^{-h}, \tau)}{2h}, \qquad h = 0.01. \tag{31}$$

The paths illustrate the fitted measure change; four paths are not a physical model validation.

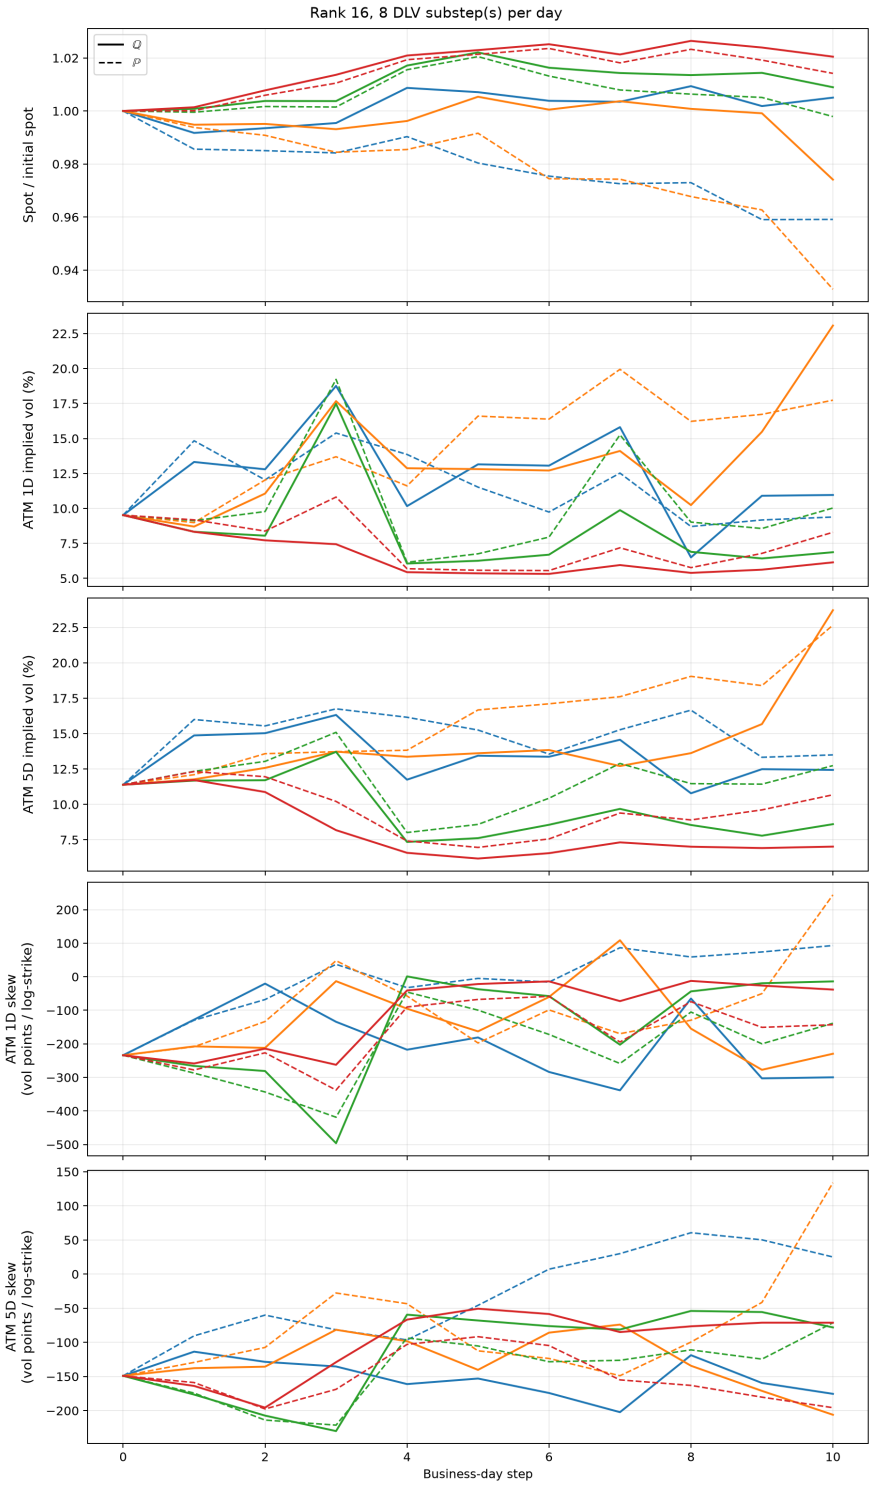


Figure 1: Matched $\mathbb{P}/\mathbb{Q}$ paths for Model D: rank 16 with eight Brownian substeps per day.

A.2 PCA Summary

The remaining figures visualize the factor systems estimated in Section 3.3. Figure 2 is the baseline risk-neutral implied-volatility surface used for all factor illustrations.

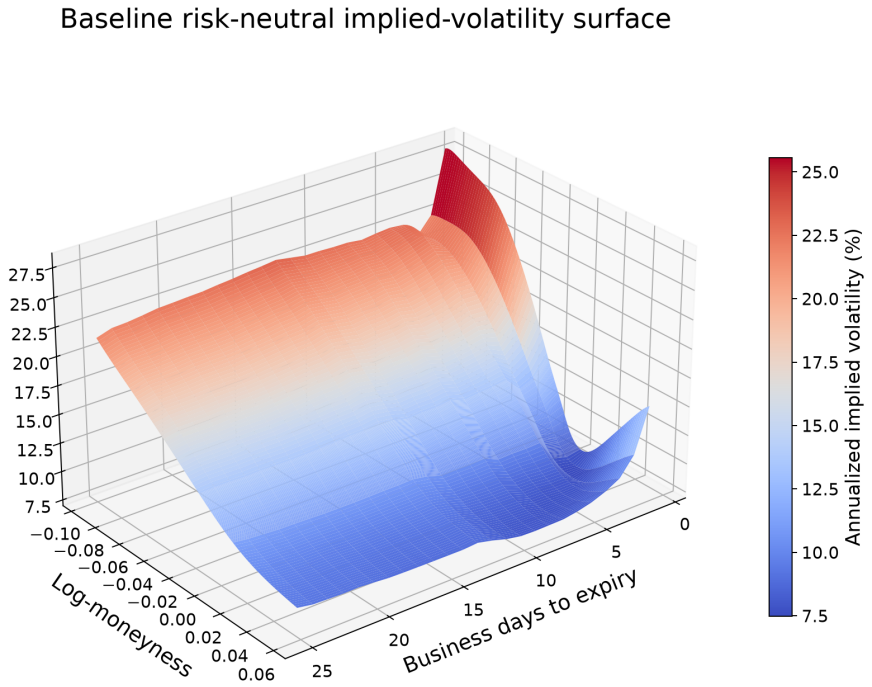


Figure 2: Baseline annualized implied volatility on the fitted log-moneyness and business-days-to-expiry grid.

Figure 3 reports individual and cumulative PCA power. The experiments retain either eight or 16 factors; the first rank above 90% is 26.

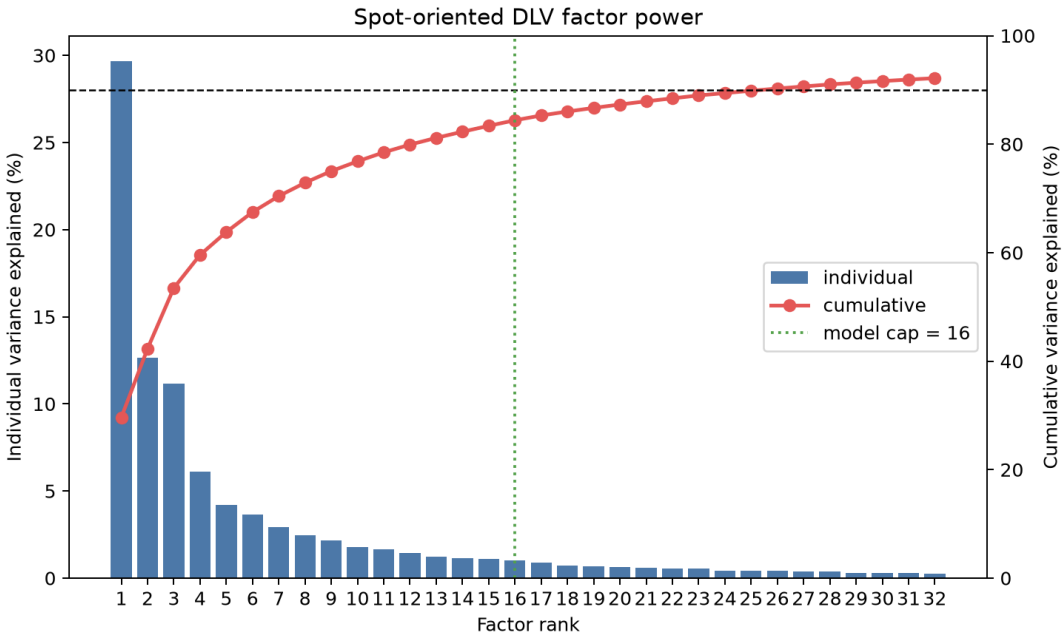


Figure 3: Individual and cumulative variance explained by the spot-oriented DLV factor system. The vertical line is the rank-16 model cap; 26 factors are required to exceed 90%.

The rank-16 fit appends PCs 9–16 to the rank-8 fit. A one-standard-deviation shock is the diffusion-only log-DLV displacement $\pm\xi_f\sqrt{1/252}$. The figures display its central implied-volatility response

$$\frac{1}{2}\left[\sigma_{\text{imp}}(\log\Sigma+\xi_f/\sqrt{252})-\sigma_{\text{imp}}(\log\Sigma-\xi_f/\sqrt{252})\right]. \tag{32}$$

These shock directions are the same under $\mathbb{P}$ and $\mathbb{Q}$: an equivalent measure change alters drift, not diffusion. Each panel deliberately retains its own vertical scale.

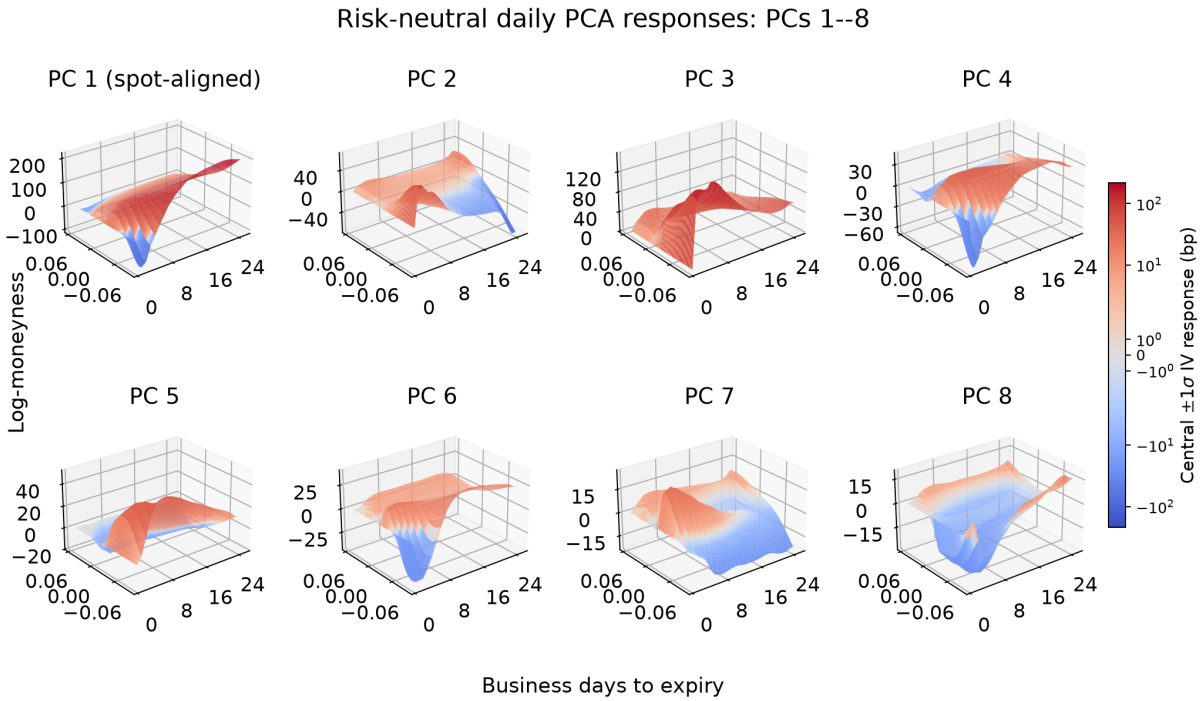


Figure 4: Risk-neutral one-standard-deviation implied-volatility responses for PCs 1–8. PC 1 is aligned to the front-spot score.

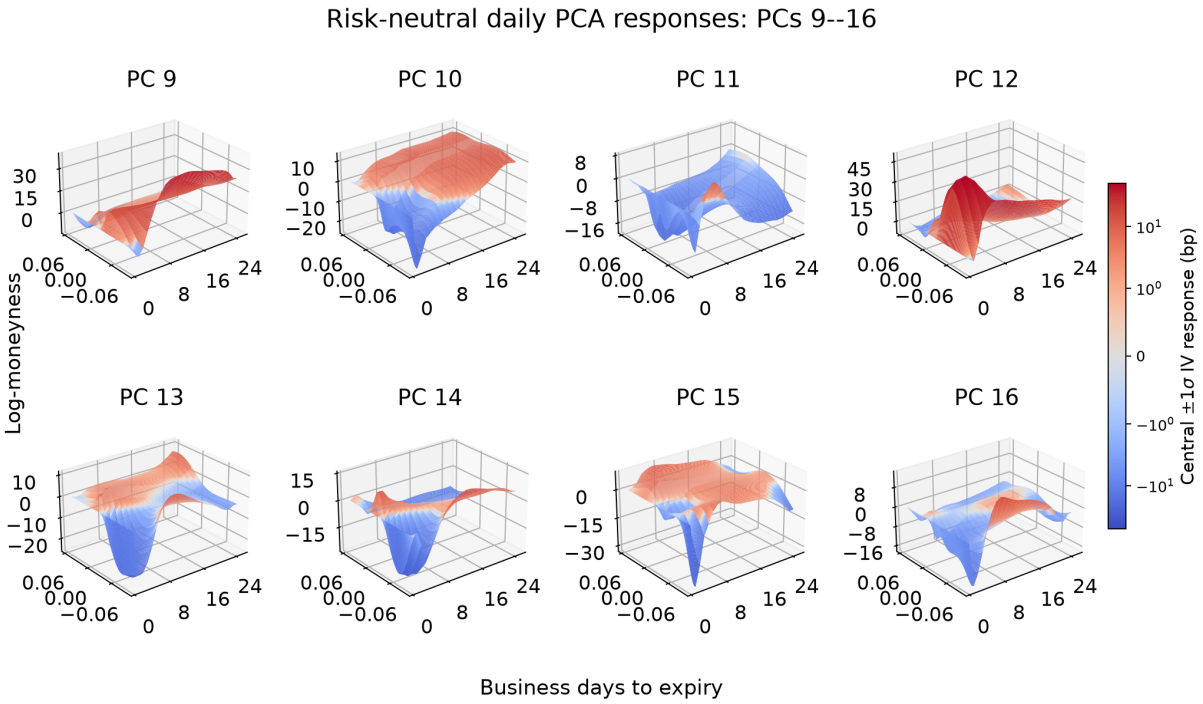


Figure 5: The additional one-standard-deviation implied-volatility responses, PCs 9–16, in the nested rank-16 model.

B Market-data and calibration pipeline

This appendix records the data and numerical pipeline behind Section 3. It follows the broader SPX pipeline described in [BHK26b], restricted here to the short-dated fixed-expiry SANOS experiment. The complete transformation is

$$\begin{aligned}\text{raw IvyDB quotes} &\longrightarrow \text{daily cleaning and parity-forward enrichment}\\ &\longrightarrow \text{arbitrage-free SANOS eligibility surface}\\ &\longrightarrow \text{European short-horizon quote selection}\\ &\longrightarrow \text{daily 25-by-200 calibration tensor}\\ &\longrightarrow \text{direct or affine end-to-end fit.}\end{aligned}$$

B.1 Raw data and cleaning

The source is OptionMetrics IvyDB US data accessed through WRDS. For each eligible SPX market date we load the index level and previous index level, option strikes and calendar expiries, call and put bid and ask quotes, volume, open interest, AM/PM settlement flags, and the zero curve. The retained sample runs from 3 January 2022 through 29 August 2025 and contains 891 valid full-day NYSE sessions.

The raw cleaning step removes known bad ticker–date observations, missing or invalid quotes, invalid settlement flags, same-day expiries, irregular expiries, and quotes whose bid does not exceed 0.011. Business days and year fractions to expiry are then computed; an AM-settled expiry is treated as 0.2 of a business day earlier than its corresponding PM expiry.

For each expiry we recompute the forward from put–call parity rather than use the delayed IvyDB forward table. With discount factor D, define the discounted-forward bounds at a strike K by

$$G^{\mathrm{ask}}(K) = C^{\mathrm{ask}}(K) - P^{\mathrm{bid}}(K) + DK, \qquad G^{\mathrm{bid}}(K) = C^{\mathrm{bid}}(K) - P^{\mathrm{ask}}(K) + DK.$$

Their midpoint is averaged across strikes with inverse-spread weights,

$$\widehat{G} = \frac{\sum_K w_K \, [G^{\mathrm{ask}}(K) + G^{\mathrm{bid}}(K)]/2}{\sum_K w_K}, \qquad w_K = \frac{1}{\max\{G^{\mathrm{ask}}(K) - G^{\mathrm{bid}}(K), 10^{-3}\}}.$$

The resulting forward is used to construct the tightest call-equivalent quote, forward-normalized strikes and prices, implied volatilities, and Greeks. An expiry must contain at least ten strikes with a traded call or put and at least four strikes with both sides traded. The normalized spread must exceed 10^{-6} and retained implied volatilities must lie between 0.01 and 5.

B.2 Static preprocessing and quote selection

Before applying the model-specific horizon, the pipeline retains quotes with vega per unit total volatility of at least 10^{-3}, log volume and log open interest of at least one, and at least ten surviving options per expiry. The nominal per-day quote cap is 100,000 and is therefore non-binding here.

For data-quality control and to construct a stable ATM curve, every daily book is passed through the linear SANOS preprocessing of [BHK$^+$26a]. The strike grid is extended, each expiry is fitted inside the quoted bid–ask intervals, negligible-density redundant strikes are removed, and the joint surface fit enforces both butterfly and calendar-spread constraints. The smoothness

parameter is 0.25. Expiries and strikes for which this arbitrage-free auxiliary fit is unavailable are not passed to the final calibration.

This auxiliary SANOS surface determines eligibility, the ATM term structure, and the floating model support. It is not the calibration target: the JAX fits in Section 3.2 use the original retained call and put bid–ask intervals. No linear-programming state or fitted midpoint is supplied as a target.

The final selection keeps European quotes no farther than 25 business days from the observation date, excludes the 0.2-business-day slice, and retains the standardized log-moneyness interval $[-2.5, 1.5]$. This leaves 2,192,458 quotes, between 1,590 and 5,096 per date.

B.3 Daily calibration tensor

The model expiry grid is

$$T_j = \frac{j}{252}, \qquad j = 1, \ldots, 25.$$

For each date, 100 interior-strike candidates are generated from the front ATM scale and 100 from the 25-day ATM scale. They are merged with a minimum log-spacing safeguard to form one 200-node support shared across all model expiries on that date. The support floats from date to date. Two boundary nodes, equal to one half of the lowest inner strike and one and a half times the highest inner strike, complete the 202-node density grid.

Each retained market quote stores its call/put flag, normalized strike, bid, ask, midpoint, business days to expiry, adjacent model-expiry indices, and linear interpolation weight. The contemporaneous SPX level is retained as a reference spot. The immutable calibration tensor pads each date to a multiple of 128 quotes and carries a Boolean mask, so padding never contributes to the objective.

At valuation time the DLV recursion produces the 25 discrete densities. The fixed 5% Black factor convolves each discrete density into a smooth call-price curve, and the two adjacent expiry curves are linearly interpolated to the market quote's exact expiry. Puts follow from put–call parity.

B.4 Optimization and retained artifacts

The direct and affine parameterizations are those in Section 3.2. Direct states are initialized from the date-specific ATM term structure and are constant across strike at initialization. Affine means use the average of these initial states; codes and loadings receive small centered Gaussian initial values. DLV squared is bounded between 10^{-4} and 9, with a differentiable barrier near the upper bound.

The loss in (16) is differentiated over every retained quote. Quote blocks of size 128 are used only to limit accelerator memory: their gradients are summed before each Adam update, so this is an exact full-sample gradient rather than stochastic minibatching. Each fit uses 300 initial epochs with a cosine learning-rate schedule from 2×10^{-2} to 10^{-3}, followed by a 700-epoch continuation from 5×10^{-3} to 5×10^{-5}. Checkpoints store either all daily direct states or the affine mean, loadings, and daily codes. The retained reports recompute every quote from the final checkpoint and produce the statistics in Table 1.

References

[BGM97] Alan Brace, Dariusz Gatarek, and Marek Musiela. The market model of interest rate dynamics. *Mathematical Finance*, 7(2):127–155, 1997.

[BHK+26a] Hans Buehler, Blanka Horvath, Anastais Kratsios, Yannick Limmer, and Raeid Saqur. Sanos - smooth arbitrage-free non-parametric option surfaces. `https://arxiv.org/abs/2601.11209`, Februar 2026.

[BHK26b] Hans Buehler, Blanka Horvath, and Anastasis Kratsios. Dysanos generative dynamic smooth arbitrage-free non-parametric option surfaces, 2026.

[BMPW22] Hans Buehler, Phillip Murray, Mikko S. Pakkanen, and Ben Wood. Deep hedging: learning to remove the drift. *Risk*, February 2022.

[BR16] Hans Buehler and Evgenii Ryskin. Discrete local volatility for large time steps (short version). Technical report, SSRN, 2016. Short version. An extended version also exists: SSRN 2642630 (2015), doi:10.2139/ssrn.2642630.

[CN09] René Carmona and Sergey Nadtochiy. Local volatility dynamic models. *Finance and Stochastics*, 13:1–48, 2009.

[CN12] René Carmona and Sergey Nadtochiy. Tangent lévy market models. *Finance and Stochastics*, 16(1):63–104, 2012.

[SW08a] Martin Schweizer and Johannes Wissel. Arbitrage-free market models for option prices: The multi-strike case. *Finance and Stochastics*, 12:469–505, 2008.

[SW08b] Martin Schweizer and Johannes Wissel. Term structures of implied volatilities: Absence of arbitrage and existence results. *Mathematical Finance*, 18:77–114, 2008.

[WBWB19] Magnus Wiese, Lianjun Bai, Ben Wood, and Hans Buehler. Deep hedging: Learning to simulate equity option markets, 2019. NeurIPS 2019 Workshop on Robust AI in Financial Services.

[Wis07] Johannes Wissel. Arbitrage-free market models for option prices. NCCR FINRISK Working Paper 428, NCCR FINRISK, ETH Zürich, 2007.

[WWP+21] Magnus Wiese, Ben Wood, Alexandre Pachoud, Ralf Korn, Hans Buehler, Phillip Murray, and Lianjun Bai. Multi-asset spot and option market simulation, 2021. arXiv:2112.06823.